


Cloud Computing Virtualization: Part 1

Hong Xu
Department of Computer Science
Spring 2020



香港城市大學
City University of Hong Kong



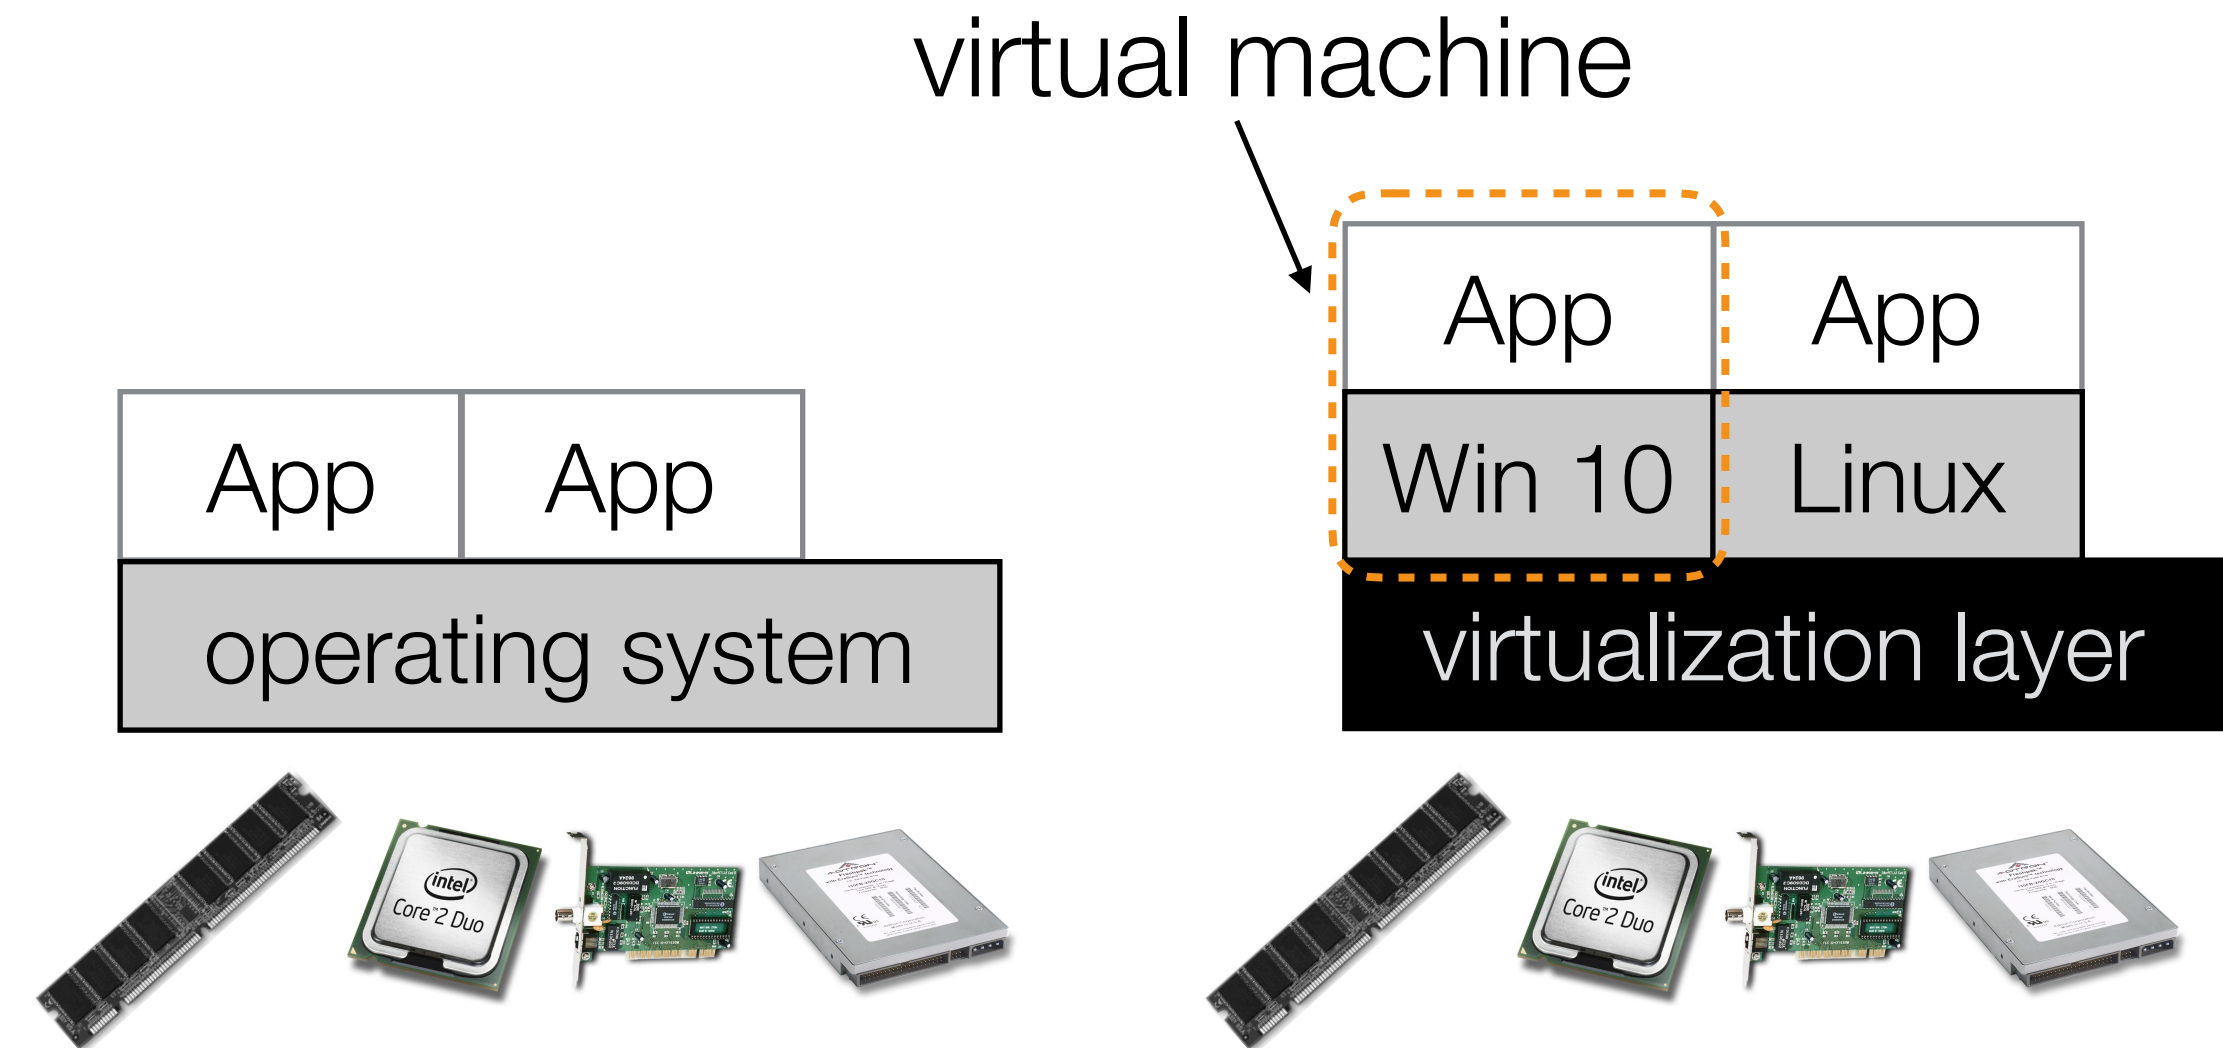
Outline

- ▶ Concepts
- ▶ Virtualization architecture
- ▶ CPU and OS basics
- ▶ Types of CPU virtualization

What is virtualization?

- ▶ Virtualization is a broad term. It can be applied to all types of resources (CPU, memory, network, etc.)
- ▶ Allows one computer to “look like” multiple computers, doing multiple jobs, by sharing the resources of a single machine across multiple environments.
 - ▶ virtualization started in 1960's in IBM's mainframe

Virtualization




‘Nonvirtualized’ system
A single OS controls all hardware platform resources

Virtualized system
It makes it possible to run multiple Virtual Machines on a single physical platform


The old model

- ▶ A server for every application
- ▶ Software and hardware are tightly coupled




The old model

- Big disadvantage: low utilization



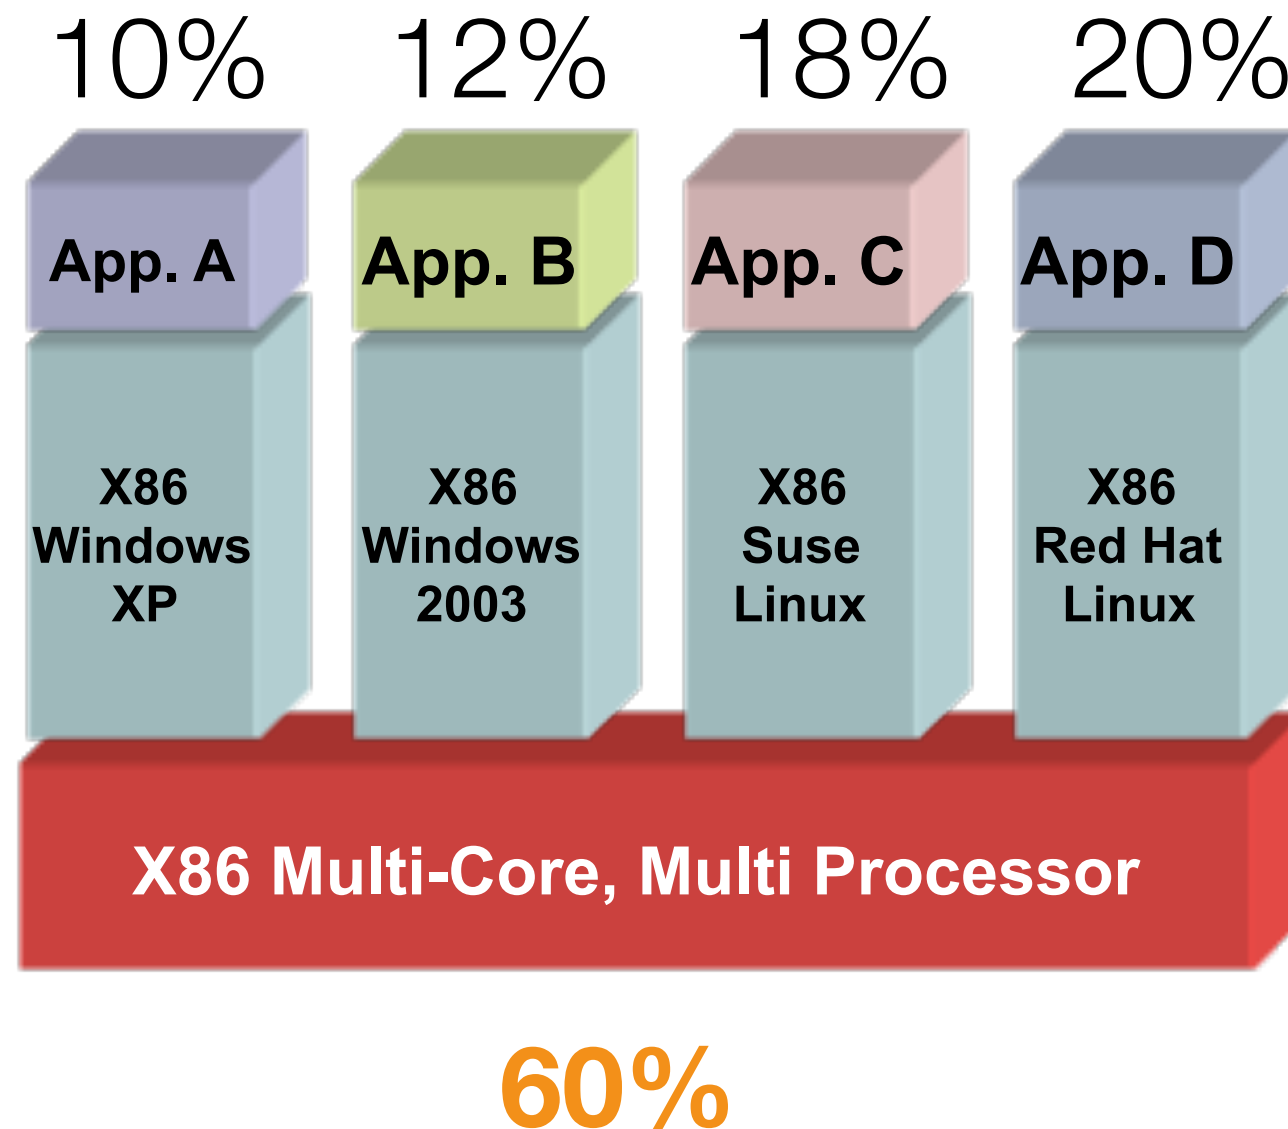
The new model

- ▶ Physical resources are virtualized. OS and applications as a single unit by encapsulating them into *virtual machines*
- ▶ Separate applications and hardware

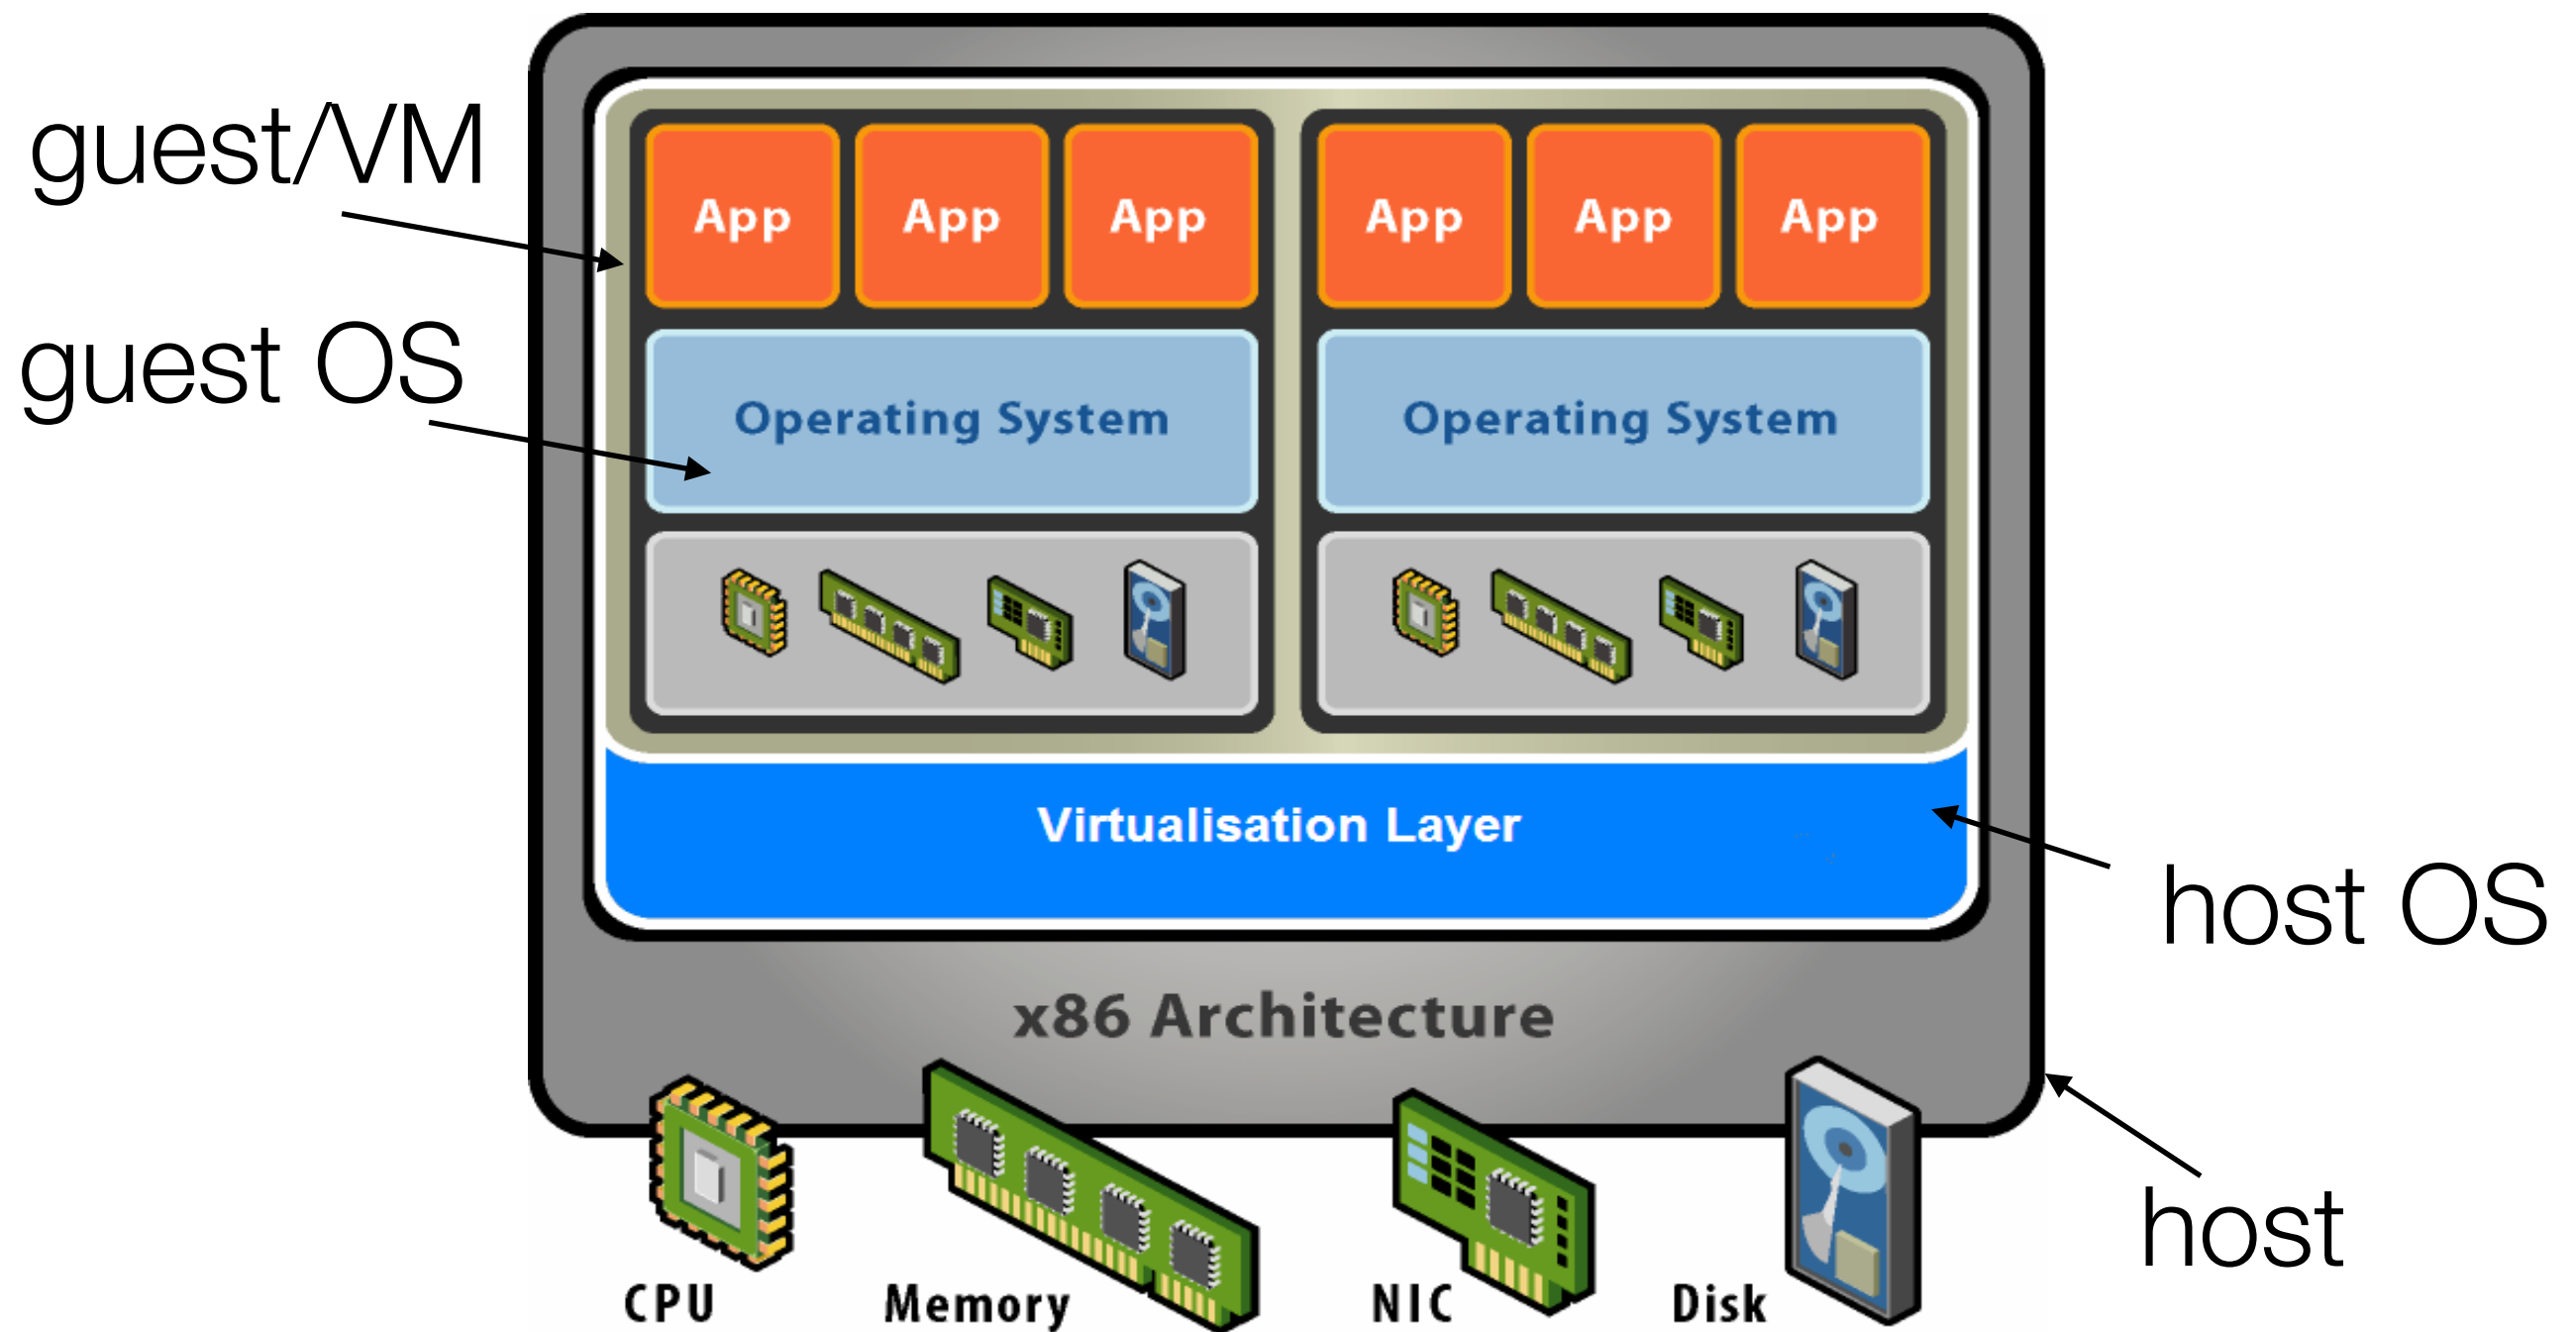


The new model

- Big advantage: improved utilization




Some terms



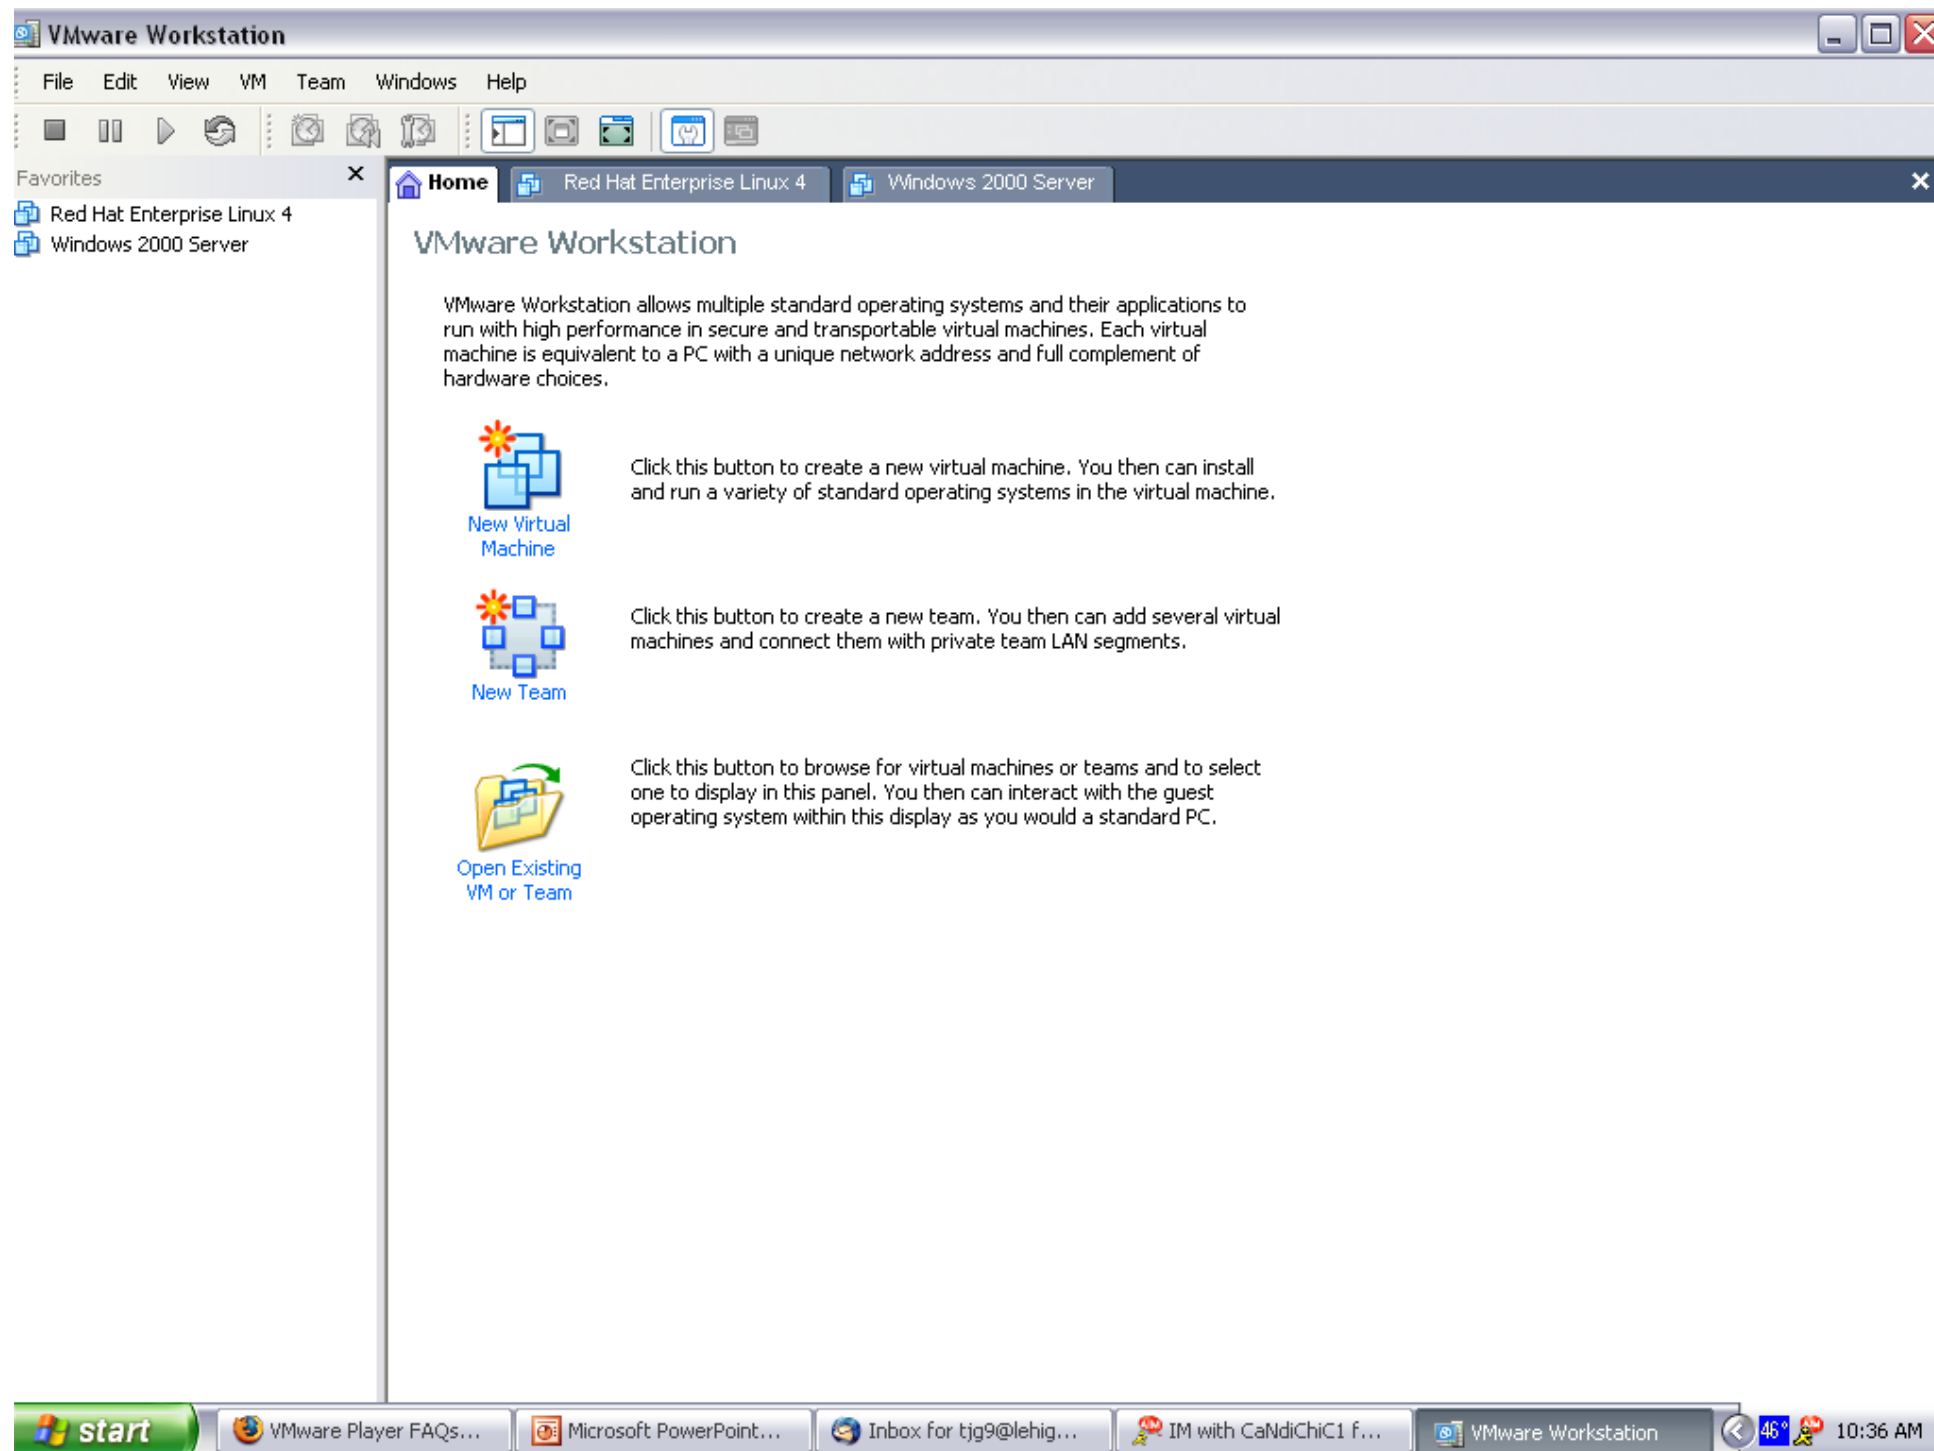
Virtualization architecture


Hosted architecture

- ▶ A hosted architecture installs and runs the virtualization layer as an application on top of an operating system



Hosted architecture example






Hosted architecture

- ▶ **Indirect** access to hardware through the host OS
 - ▶ performance penalty, usually for desktops and personal use

Hypervisor architecture

- ▶ The hypervisor architecture installs the virtualization layer, called **hypervisor**, directly on a clean x86-based system
- ▶ Installer is usually an ISO installing a tailor-made OS



Hypervisor architecture

- ▶ It has direct access to hardware resources
 - ▶ A hypervisor is more efficient than a hosted architecture and delivers greater scalability, robustness, and performance
- ▶ For production use


CPU virtualization

OS review

Credit: Prof. John Kubiatowicz's slides for
CS162, Spring 2015, UC Berkeley

What is an operating system?

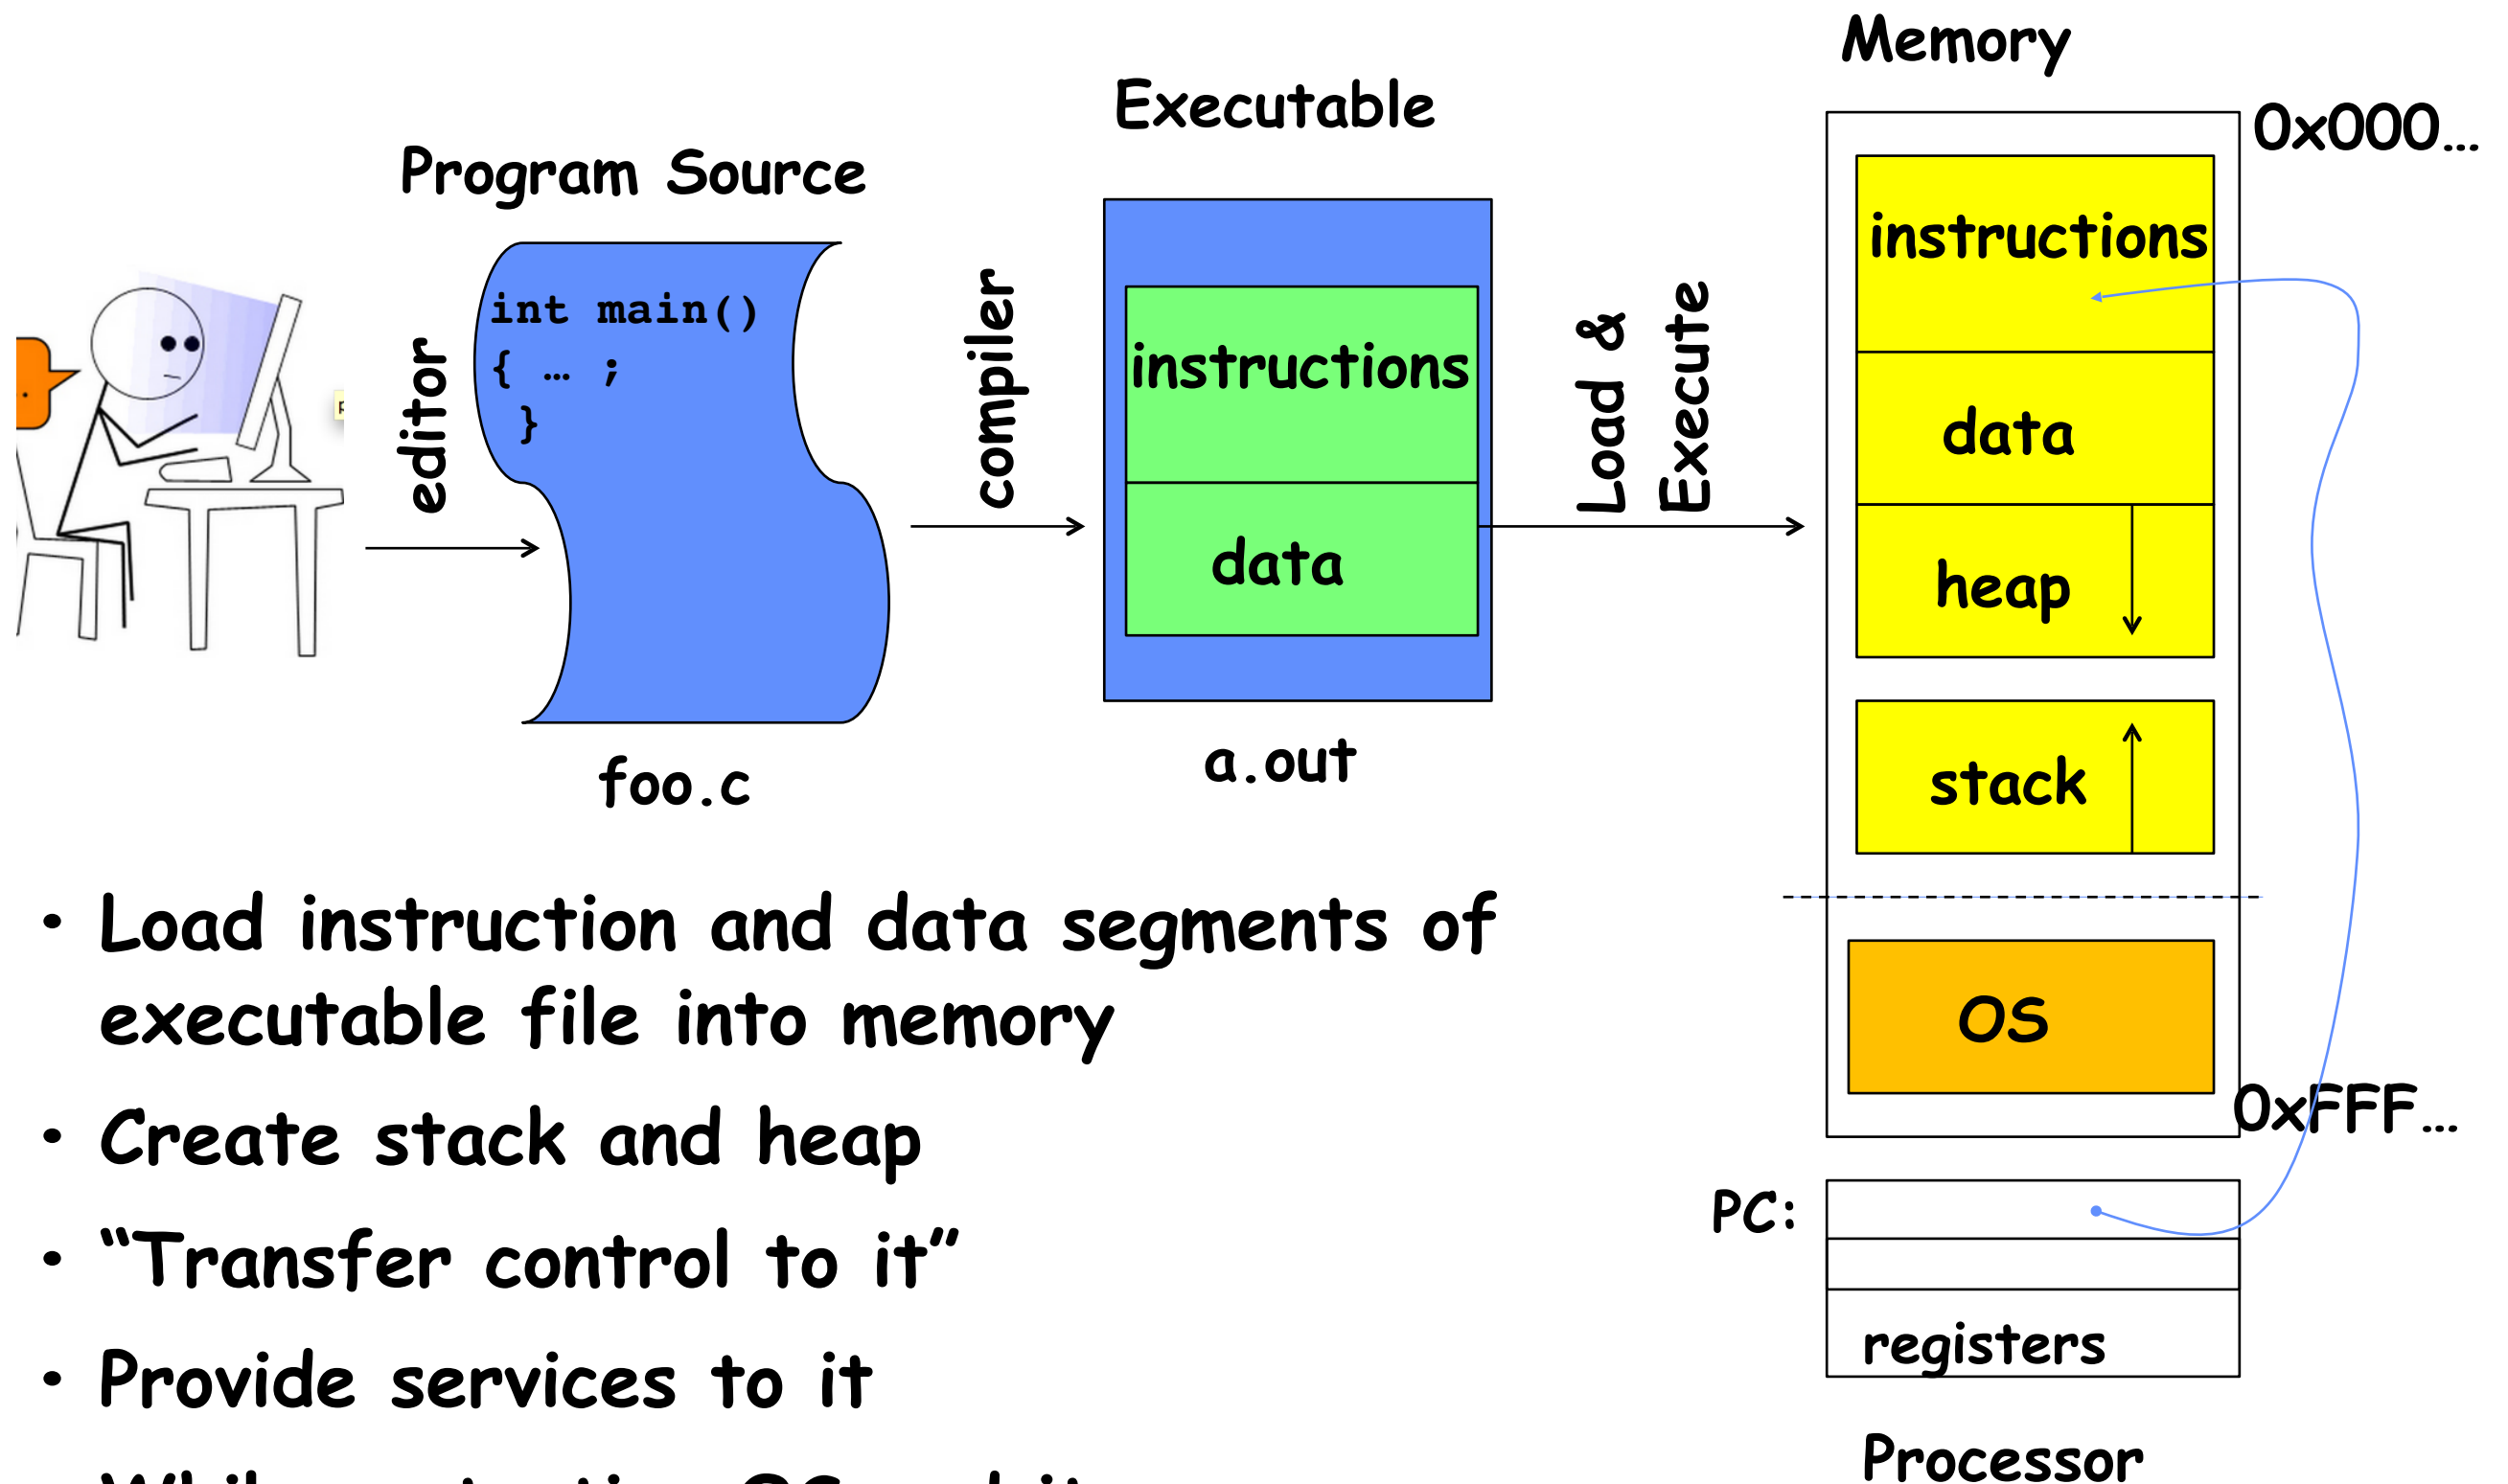
- Special layer of software that provides application access to hardware resources
 - Convenient abstraction of complex hardware devices
 - Protected access to shared resources
 - Security and authentication
 - Communication amongst logical entities



Four fundamental OS concepts

- Thread
 - Single unique execution context
 - Program Counter, Registers, Execution Flags, Stack
- Address Space w/ Translation
 - Programs execute in an address space that is distinct from the memory space of the physical machine
- Process
 - An instance of an executing program is a process consisting of an address space and one or more threads of control
- Dual Mode operation/Protection
 - Only the "system" has the ability to access certain resources
 - The OS and the hardware are protected from user programs and user programs are isolated from one another by controlling the translation from program virtual addresses to machine physical addresses

OS Bottom Line: Run Programs



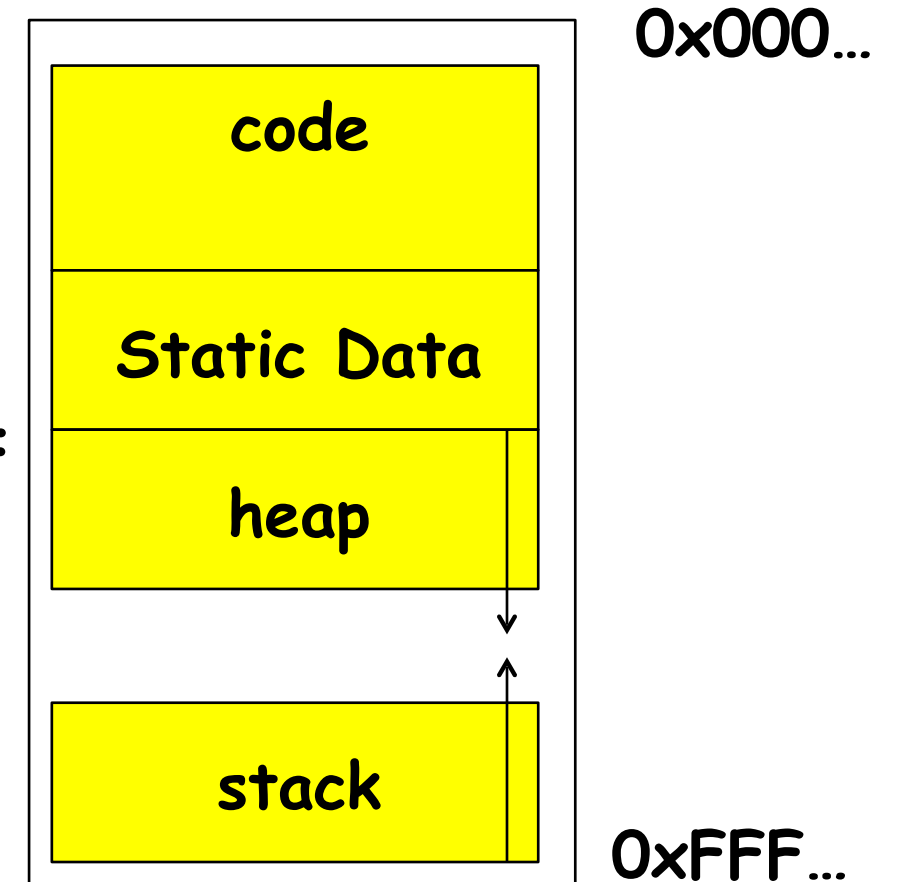
- Load instruction and data segments of executable file into memory
- Create stack and heap
- "Transfer control to it"
- Provide services to it
- While protecting OS and it

First OS Concept: Thread of Control

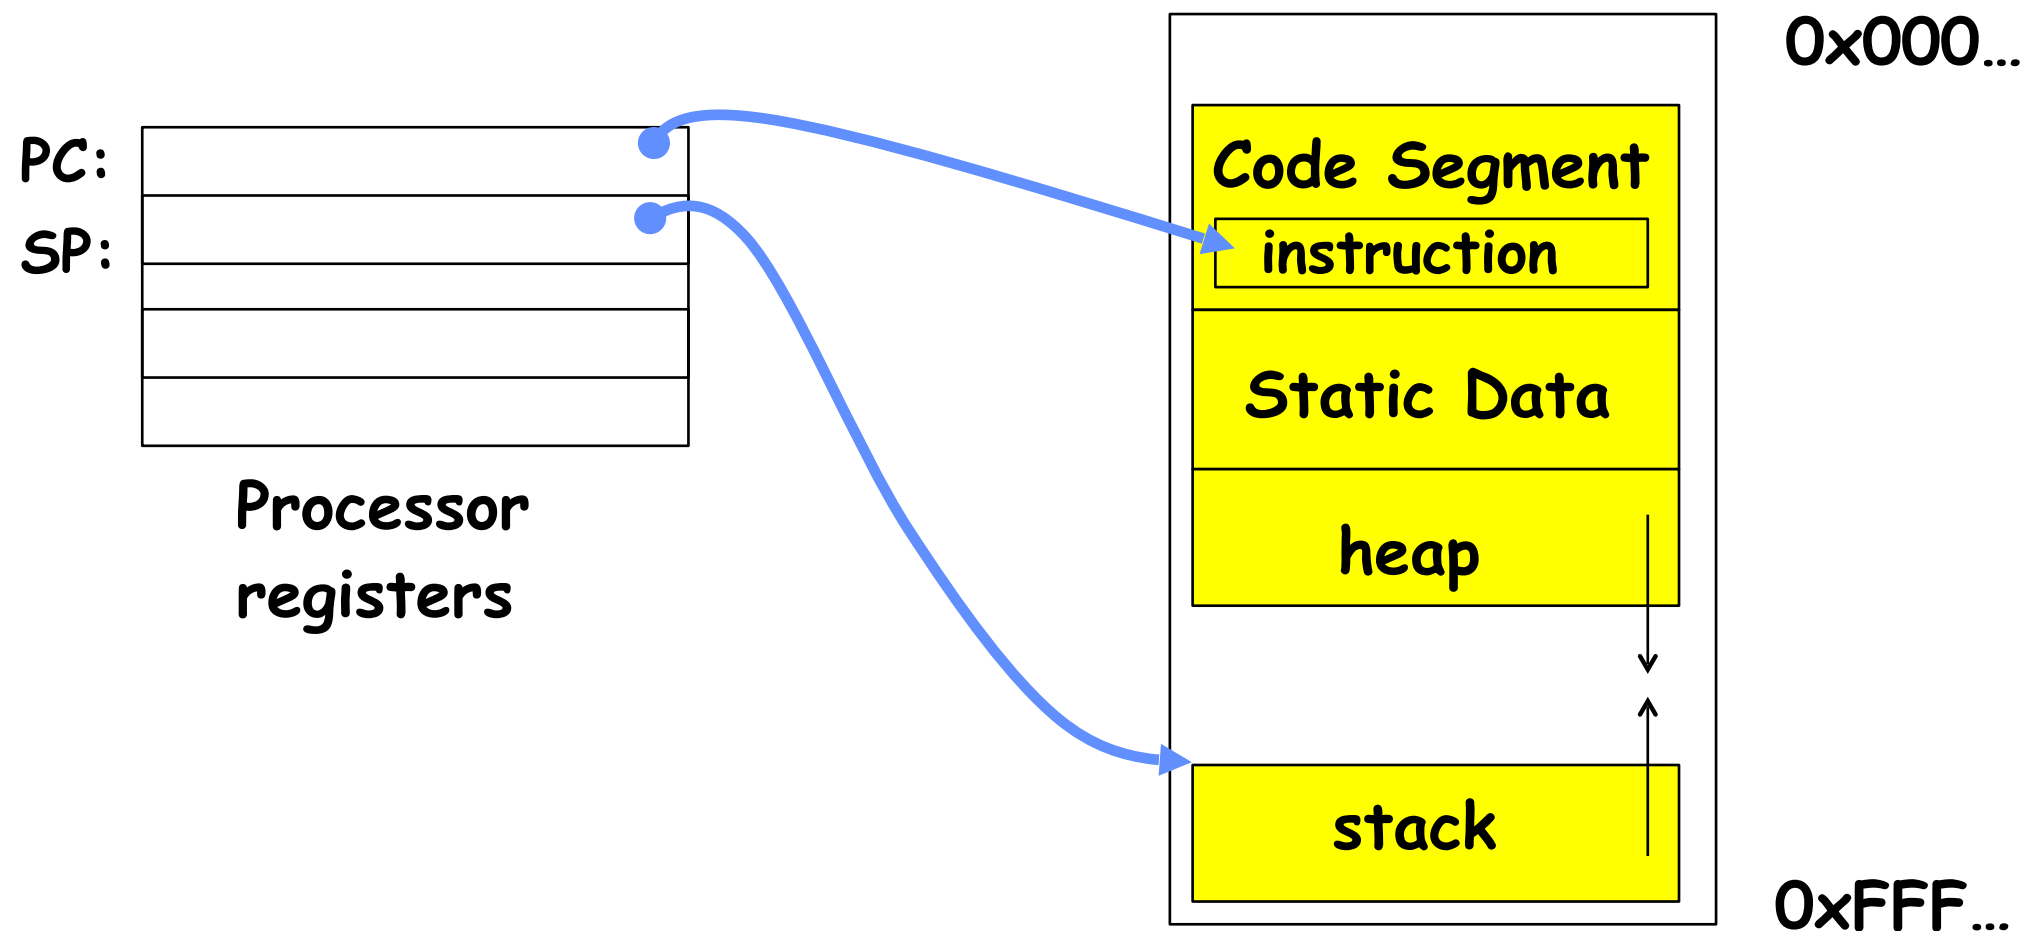
- **Thread: Single unique execution context**
 - Program Counter, Registers, Execution Flags, Stack
- A thread is executing on a processor when it is resident in the processor registers.
- PC register holds the address of executing instruction in the thread.
- Certain registers hold the context of thread
 - Stack pointer holds the address of the top of stack
 - » Other conventions: Frame Pointer, Heap Pointer, Data
 - May be defined by the instruction set architecture or by compiler conventions
- Registers hold the root state of the thread.
 - The rest is "in memory"

Second OS Concept: Program's Address Space

- **Address space \Rightarrow the set of accessible addresses + state associated with them:**
 - For a 32-bit processor there are $2^{32} = 4$ billion addresses
- What happens when you read or write to an address?
 - Perhaps Nothing
 - Perhaps acts like regular memory
 - Perhaps ignores writes
 - Perhaps causes I/O operation
 - » (Memory-mapped I/O)
 - Perhaps causes exception (fault)

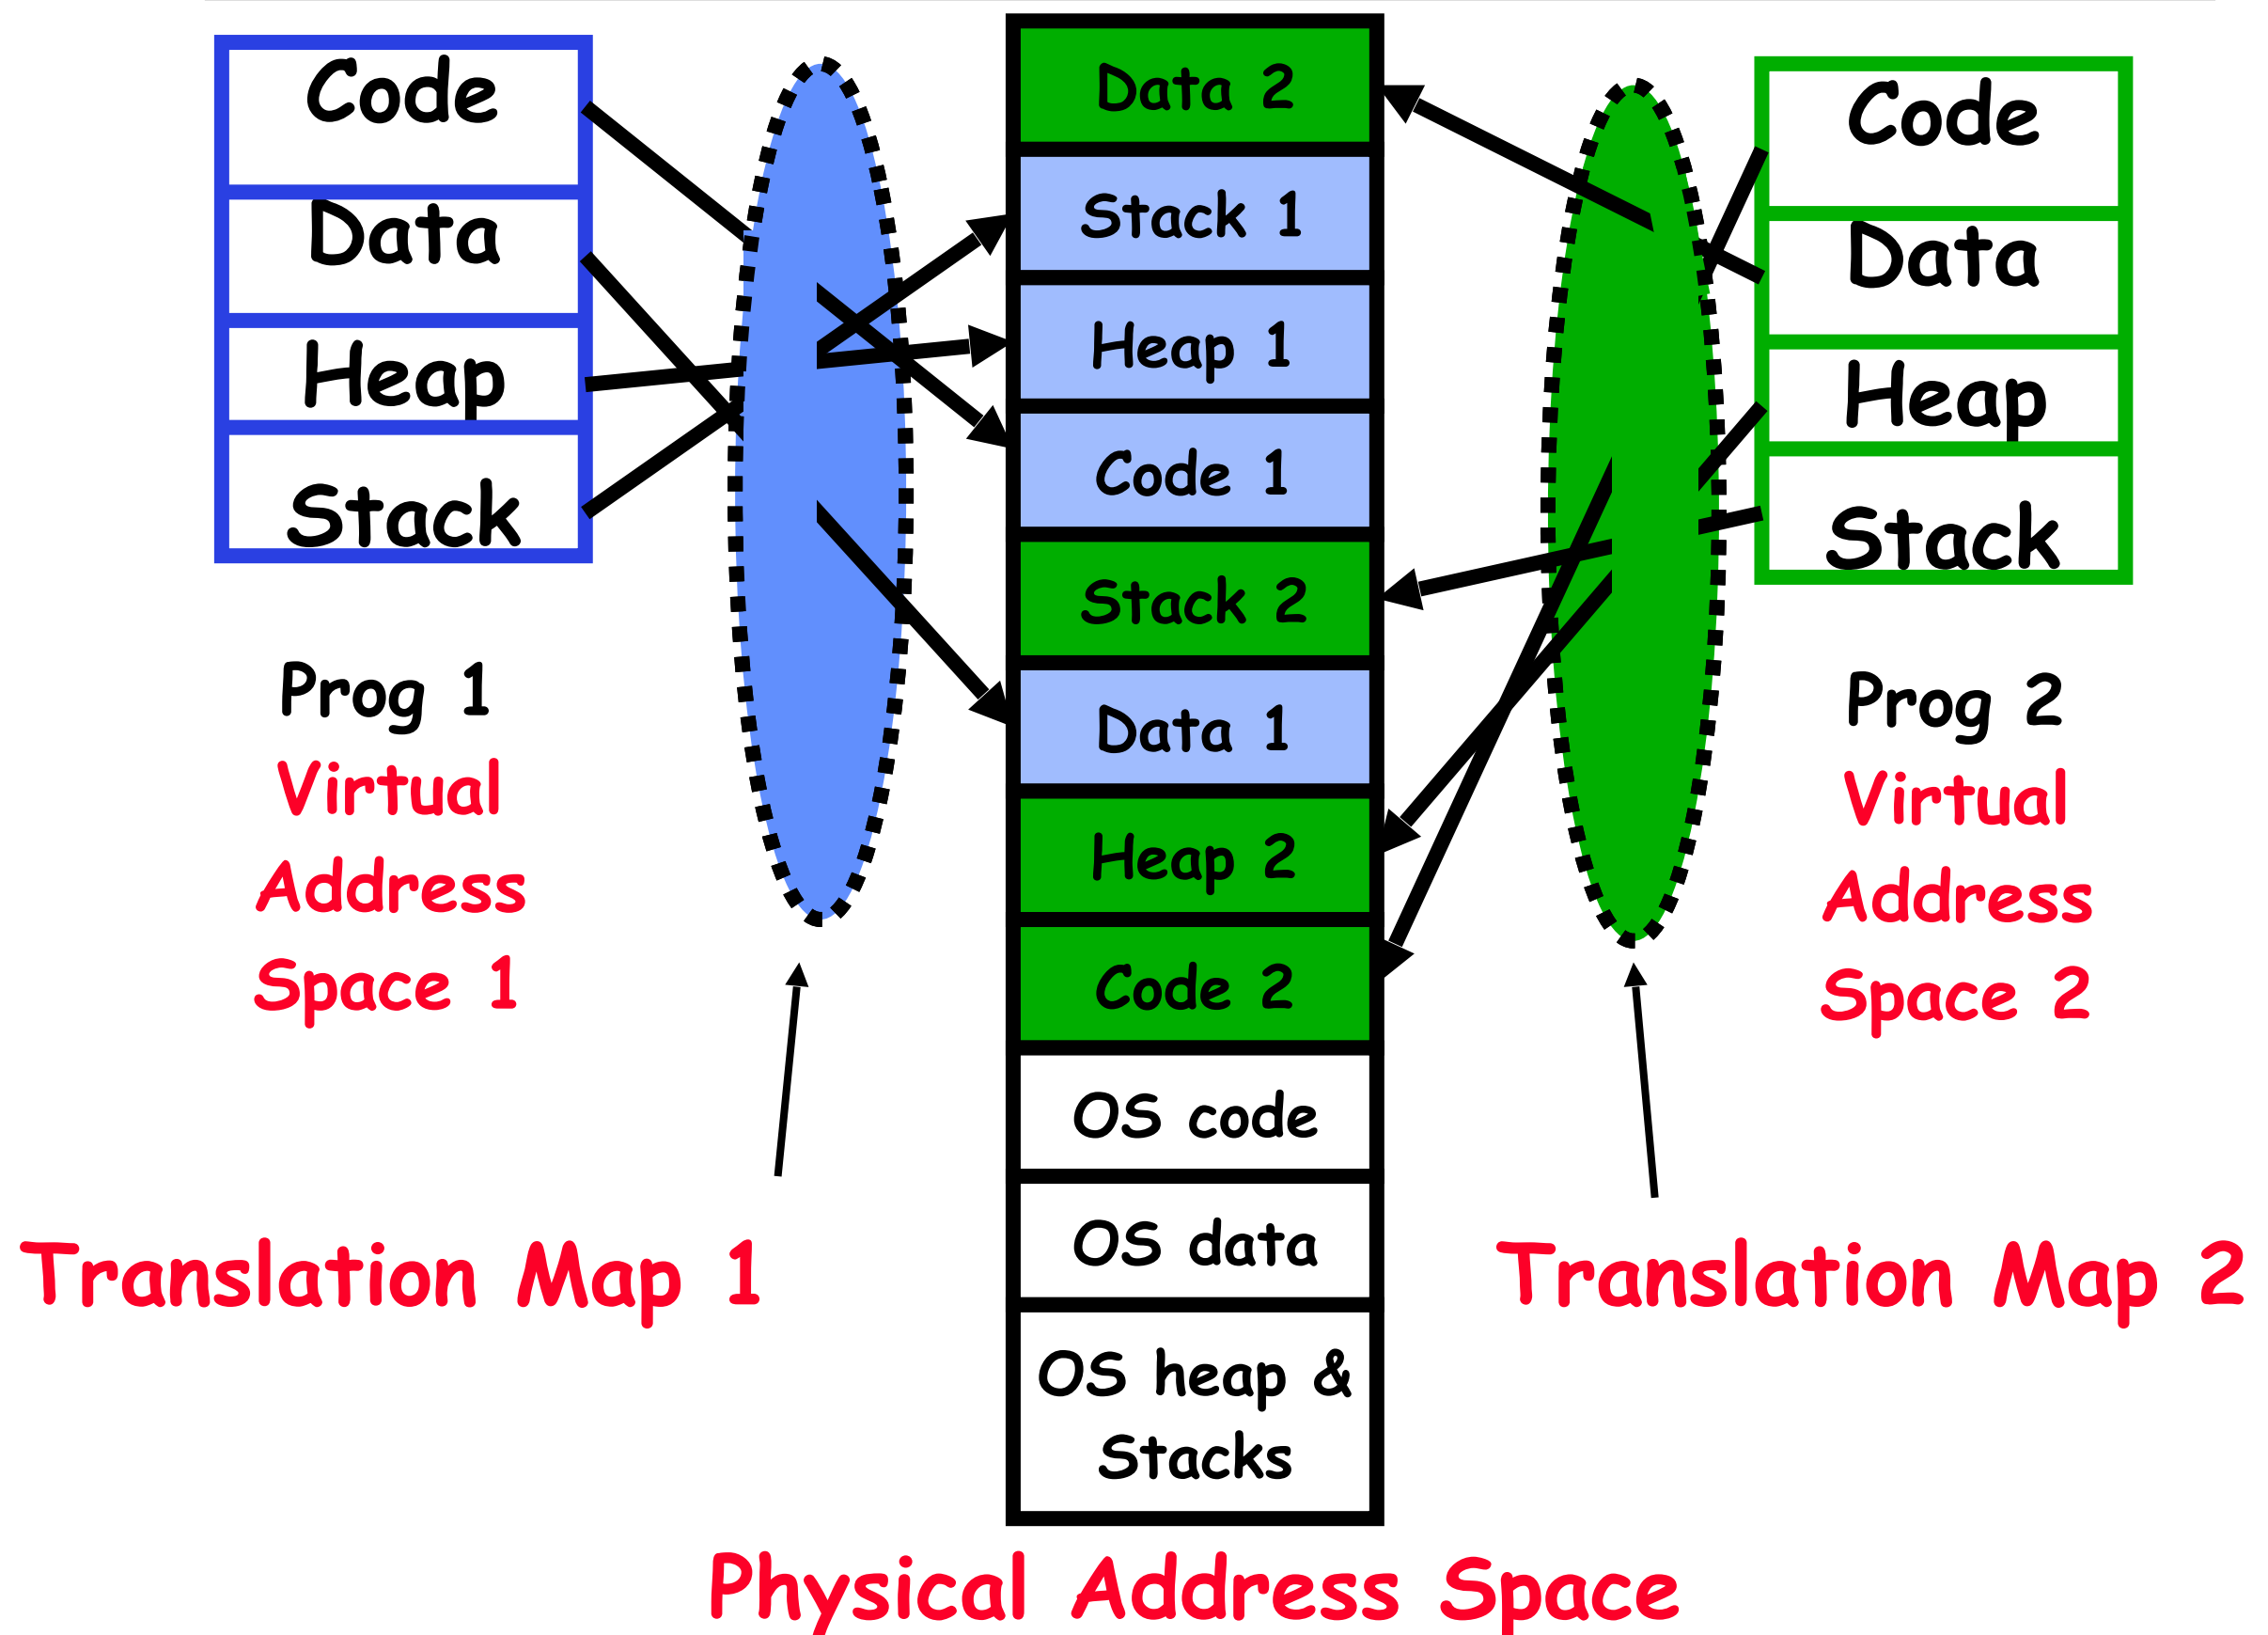


Address Space: In a Picture



- What's in the code segment? Data? (global var)
- What's in the stack segment?
 - automatic variables, register values, etc.
- What's in the heap segment?
 - variables from dynamic memory allocation (malloc, etc.)

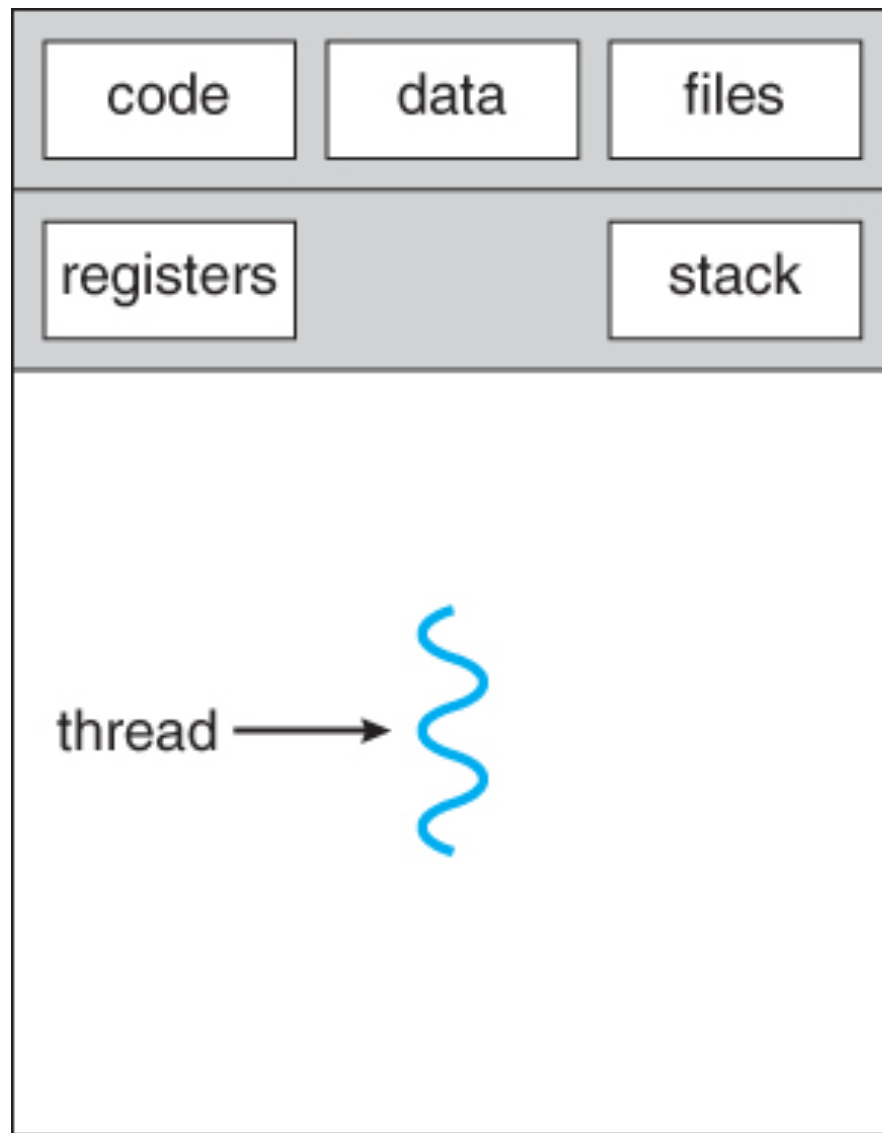
Providing Illusion of Separate Address Space:
Load new Translation Map on Switch



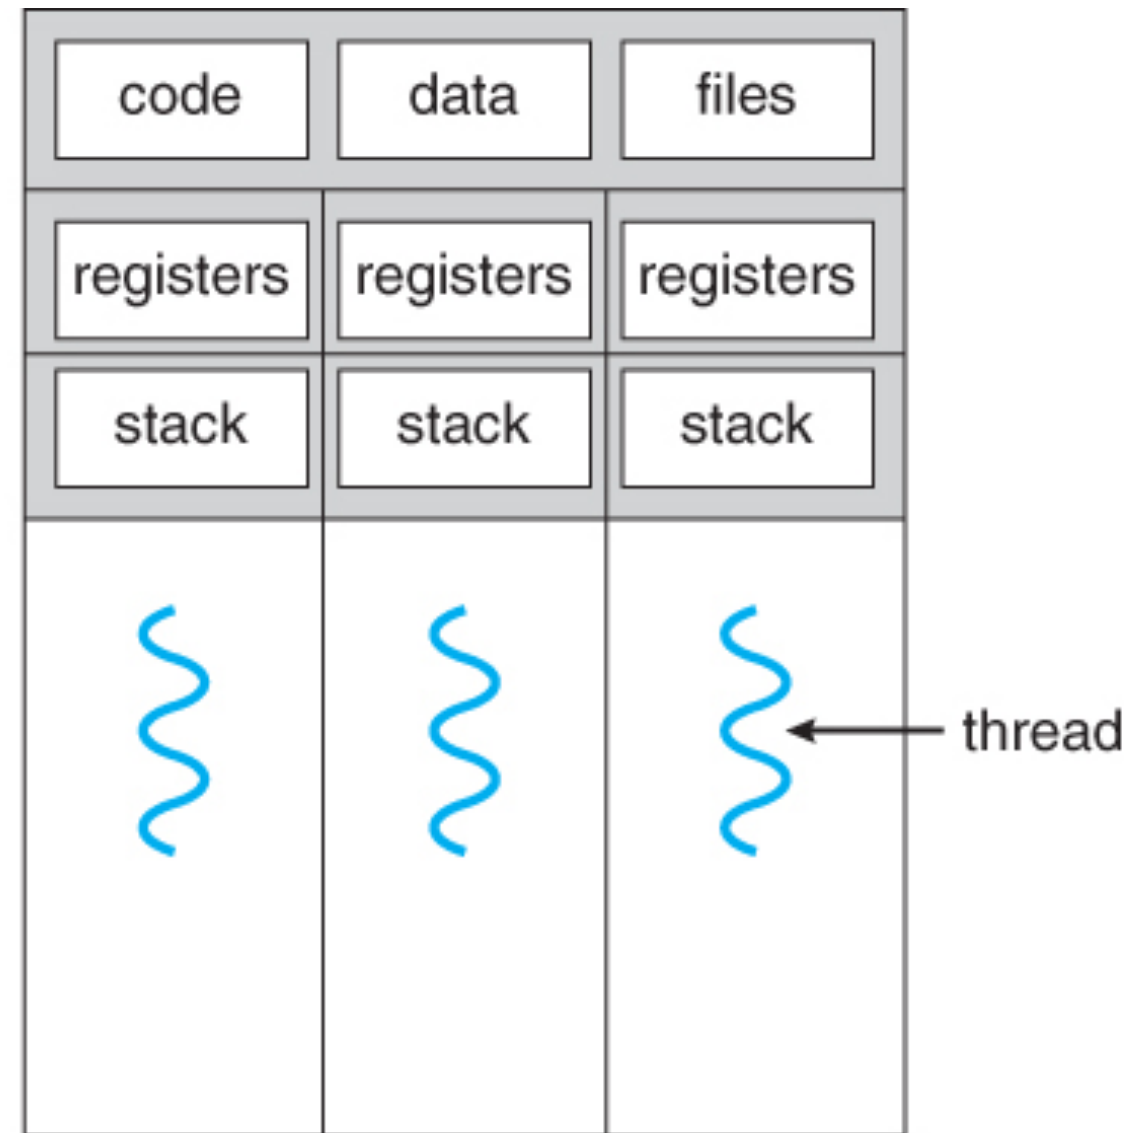
Third OS Concept: Process

- **Process: execution environment with Restricted Rights**
 - **Address Space with One or More Threads**
 - Owns memory (address space)
 - Owns file descriptors, file system context, ...
 - Encapsulate one or more threads sharing process resources
- **Why processes?**
 - **Protected from each other!**
 - **OS Protected from them**
 - Navigate fundamental tradeoff between protection and efficiency
 - Processes provides memory protection
- **Application instance consists of one or more processes**

Process vs thread



single-threaded process



multithreaded process

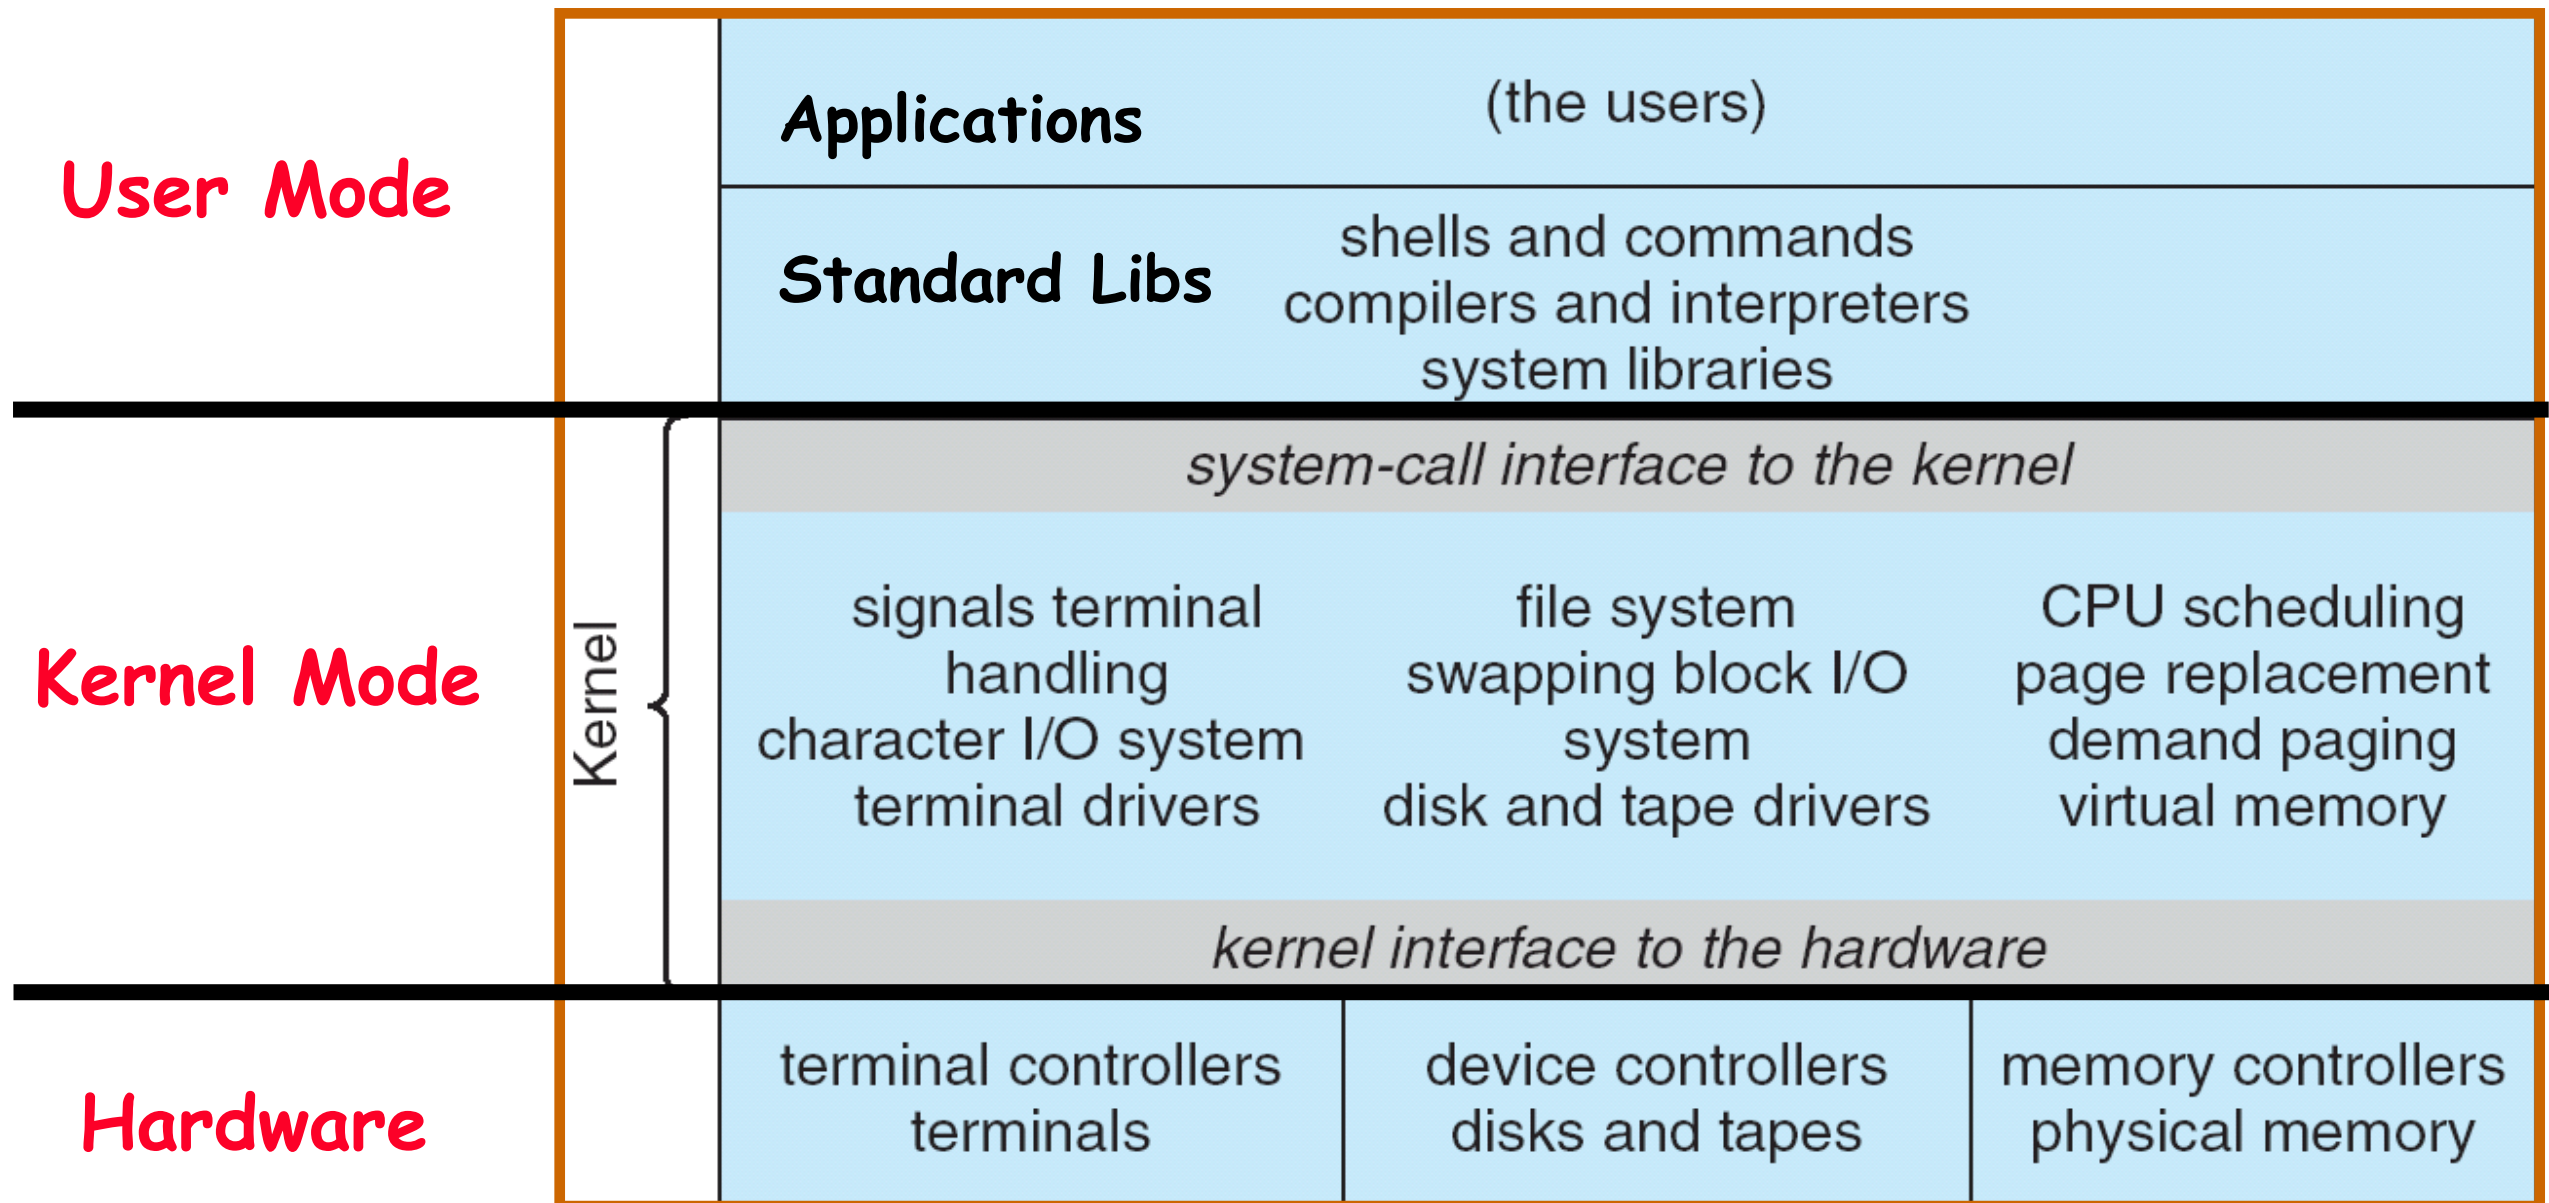
Protection

- Operating System must protect itself from user programs
 - Reliability: compromising the operating system generally causes it to crash
 - Security: limit the scope of what processes can do
 - Privacy: limit each process to the data it is permitted to access
 - Fairness: each should be limited to its appropriate share
- It must protect User programs from one another
- Primary Mechanism: limit the translation from program address space to physical memory space
 - Can only touch what is mapped in
- Additional Mechanisms:
 - Privileged instructions, in/out instructions, special registers
 - syscall processing, subsystem implementation
 - » (e.g., file access rights, etc)

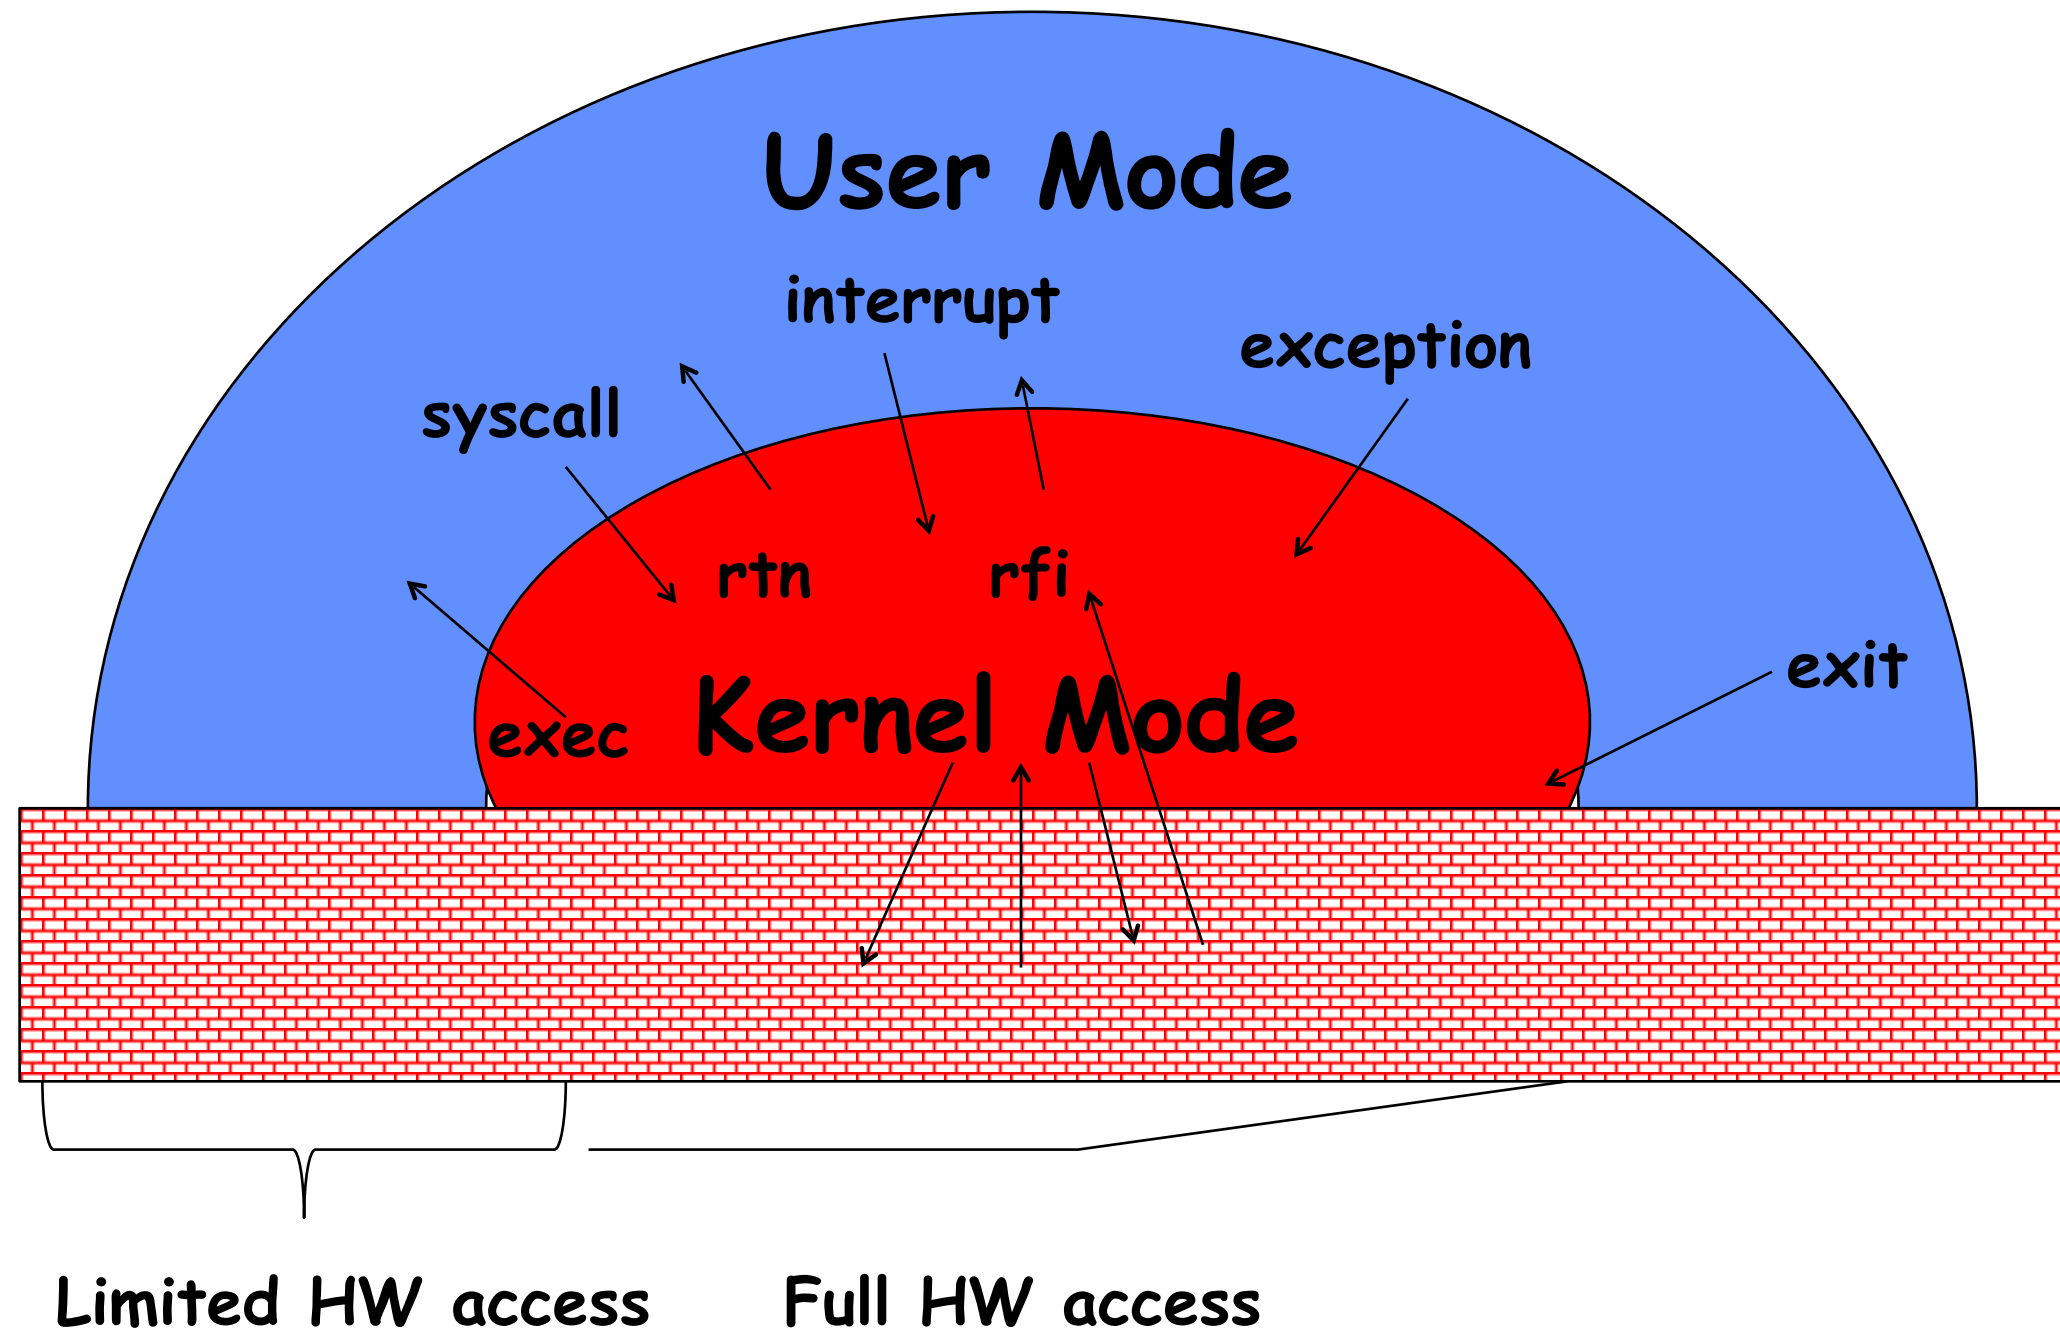
Fourth OS Concept: Dual Mode Operation

- **Hardware** provides at least two modes:
 - “Kernel” mode (or “supervisor” or “protected”)
 - “User” mode: Normal programs executed
- What is needed in the hardware to support “dual mode” operation?
 - a bit of state (user/system mode bit)
 - Certain operations / actions only permitted in system/kernel mode
 - » In user mode they fail or trap
 - User- > Kernel transition sets system mode AND saves the user PC
 - » Operating system code carefully puts aside user state then performs the necessary operations
 - Kernel- > User transition clears system mode AND restores appropriate user PC
 - » return-from-interrupt

For example: UNIX System Structure

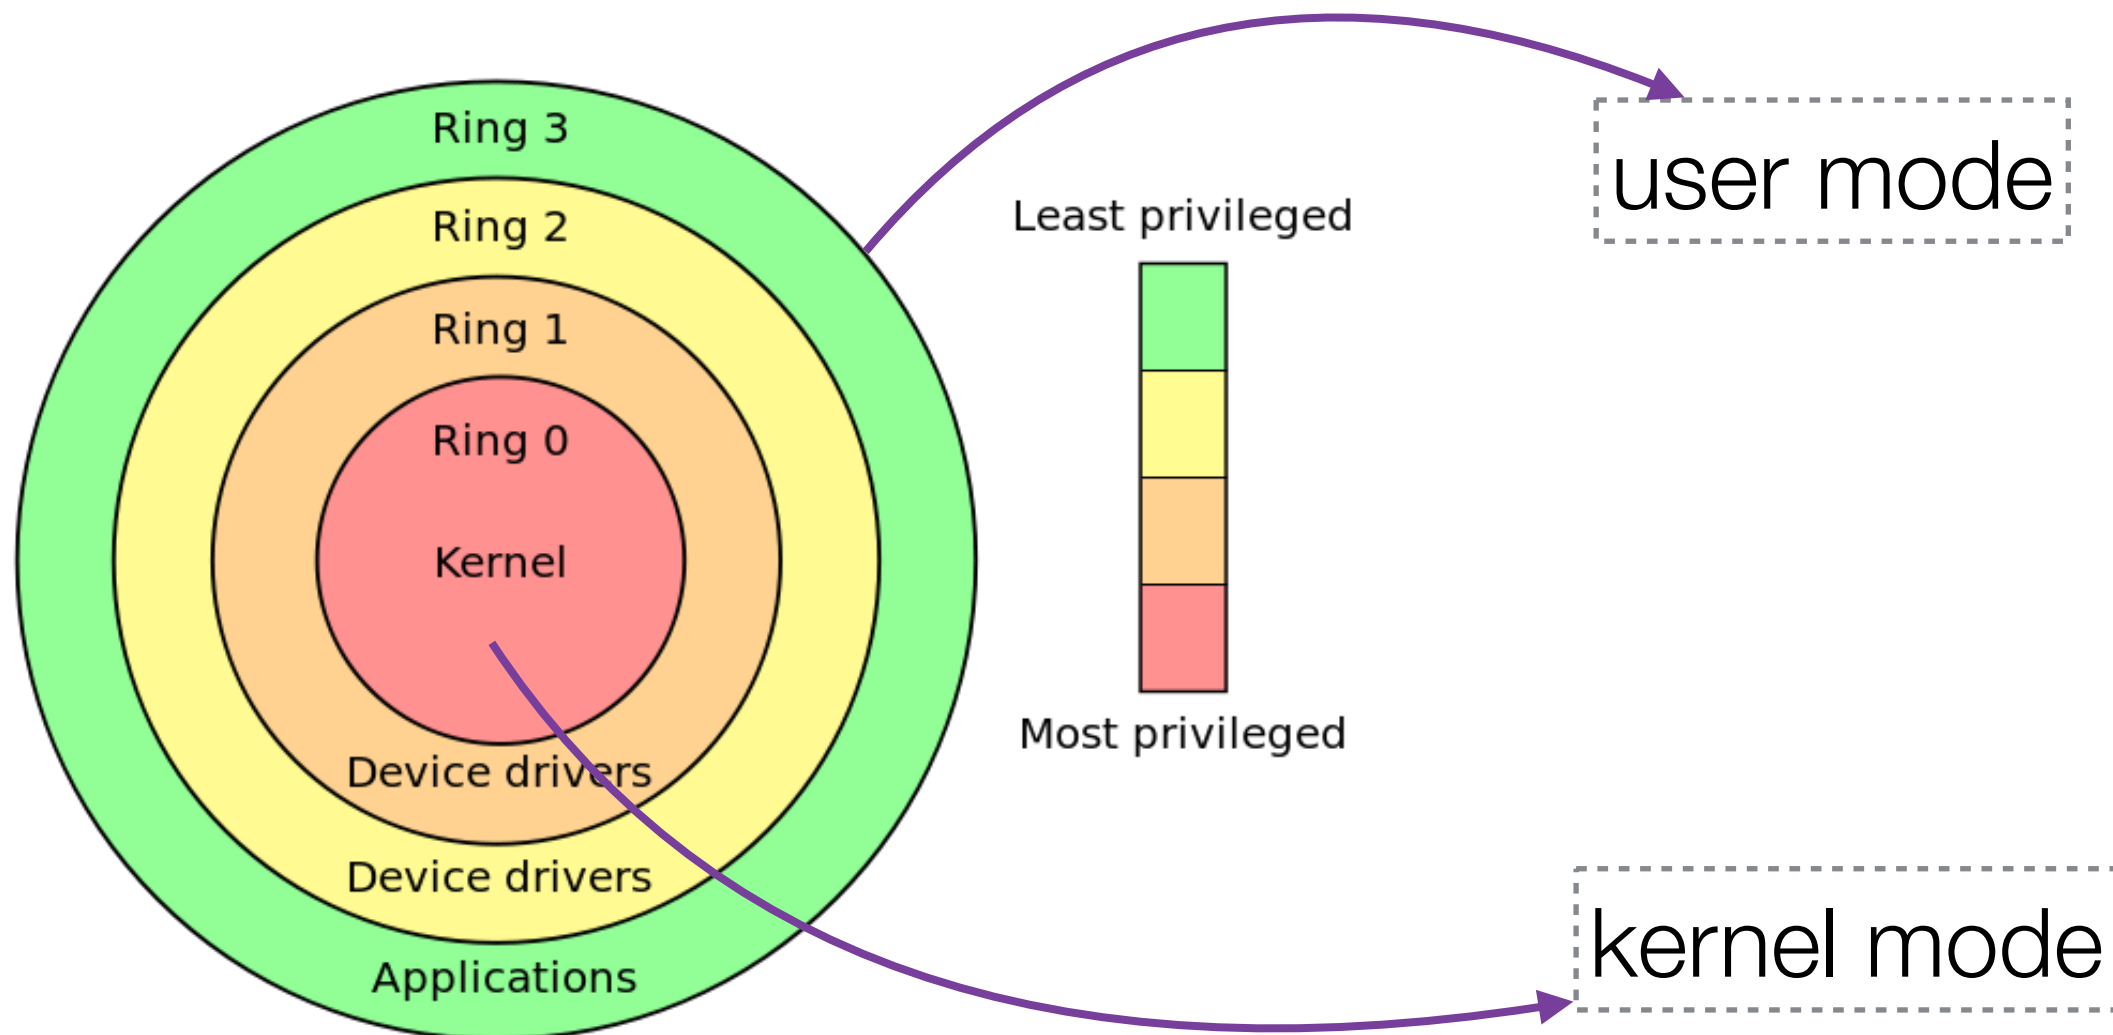


User/Kernel (Privileged) Mode



Protection rings

- ▶ Enforced in hardware in x86 architectures



Source: Wikipedia

How to virtualize a CPU?

- ▶ Basically, the CPU does not care whether you are the guest OS or not
 - ▶ Have the PC point to somewhere in the RAM
- ▶ If it's unprivileged code, code that execute in userspace
 - ▶ It's safe to run no matter it is from the guest OS or host OS
 - ▶ Why?

What about privileged code?

- ▶ Currently, there're 3 implementations
 - ▶ Full virtualization
 - ▶ Para virtualization
 - ▶ Hardware-assisted virtualization

Types of CPU virtualization

Full virtualization

Ring 3	Guest applications
Ring 2	
Ring 1	Guest OS kernel
Ring 0	Hypervisor, Host OS
Host hardware	

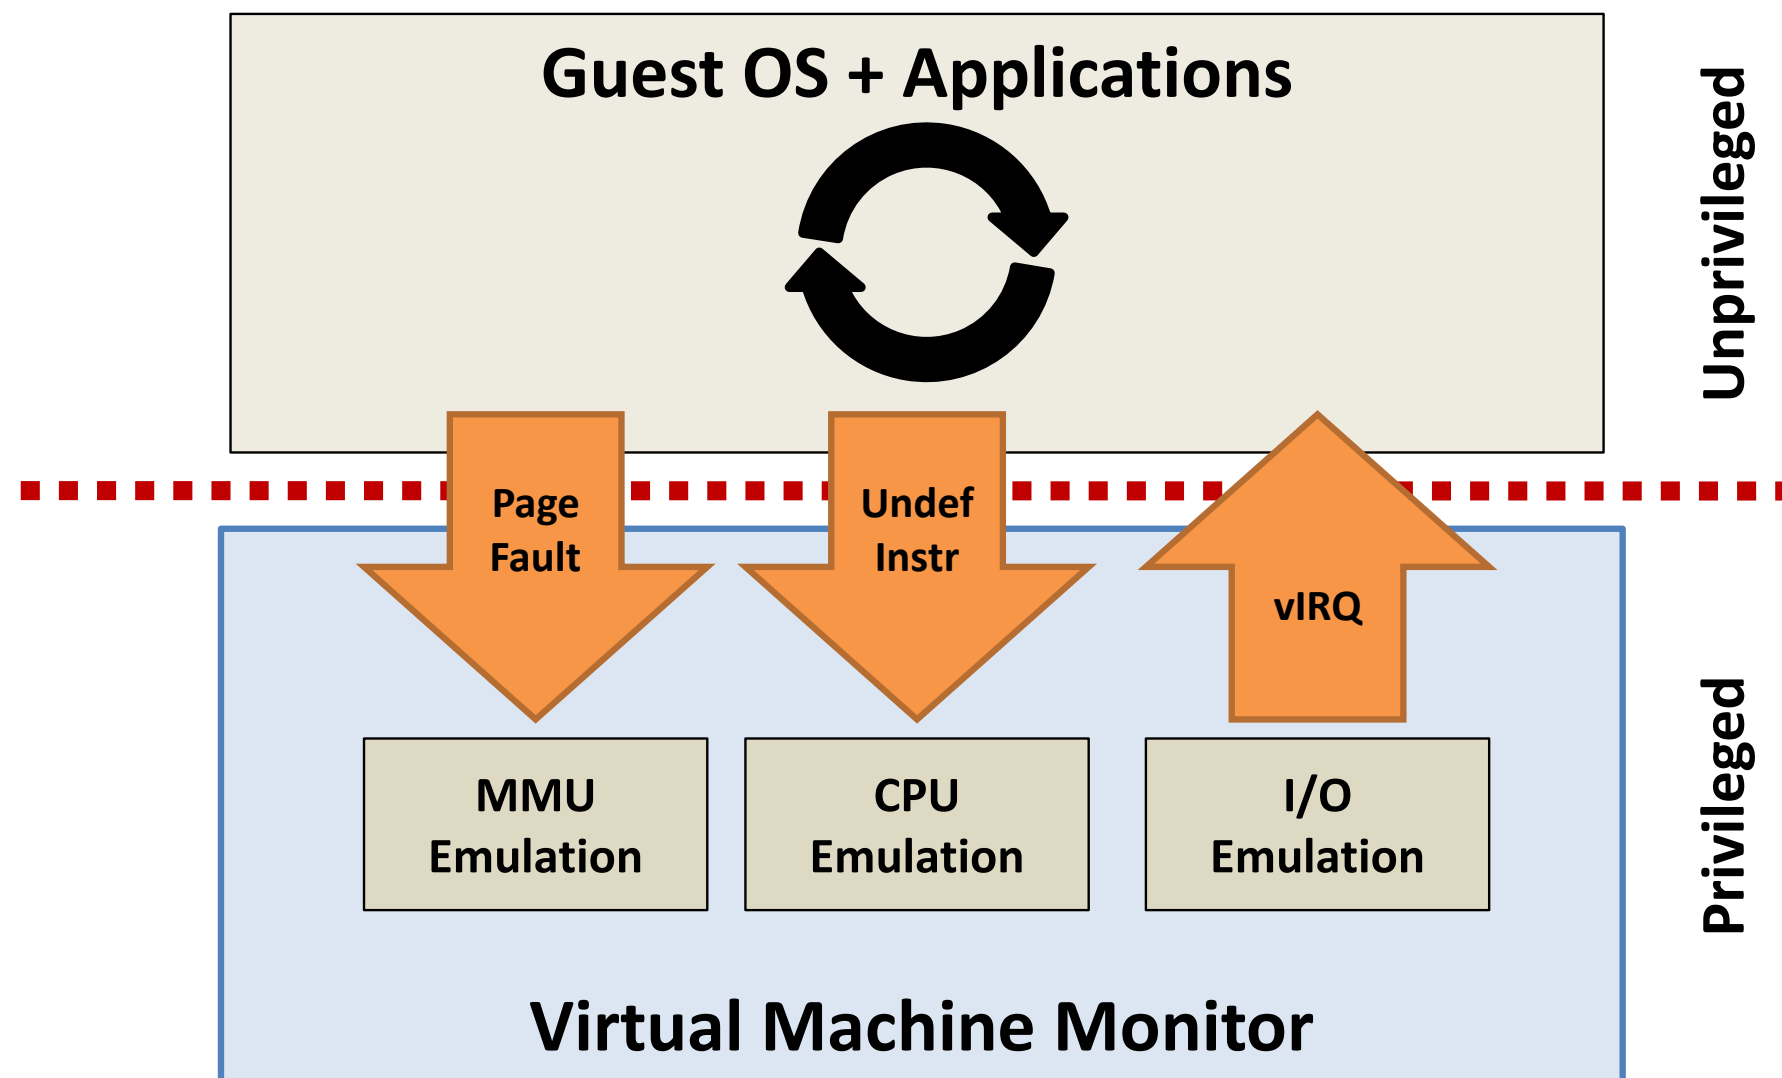


hardware emulation

binary translation

Full virtualization

- ▶ Hardware is emulated by the hypervisor



Full virtualization

- ▶ The hypervisor presents a complete set of emulated hardware to the VM's guest operating system, including the CPU, motherboard, memory, disk, disk controller, and network cards.
- ▶ For example, Microsoft Virtual Server 2005 emulates an Intel 21140 NIC card and Intel 440BX chipset.
- ▶ Regardless of the actual physical hardware on the host system, the emulated hardware remains the same.

Full virtualization

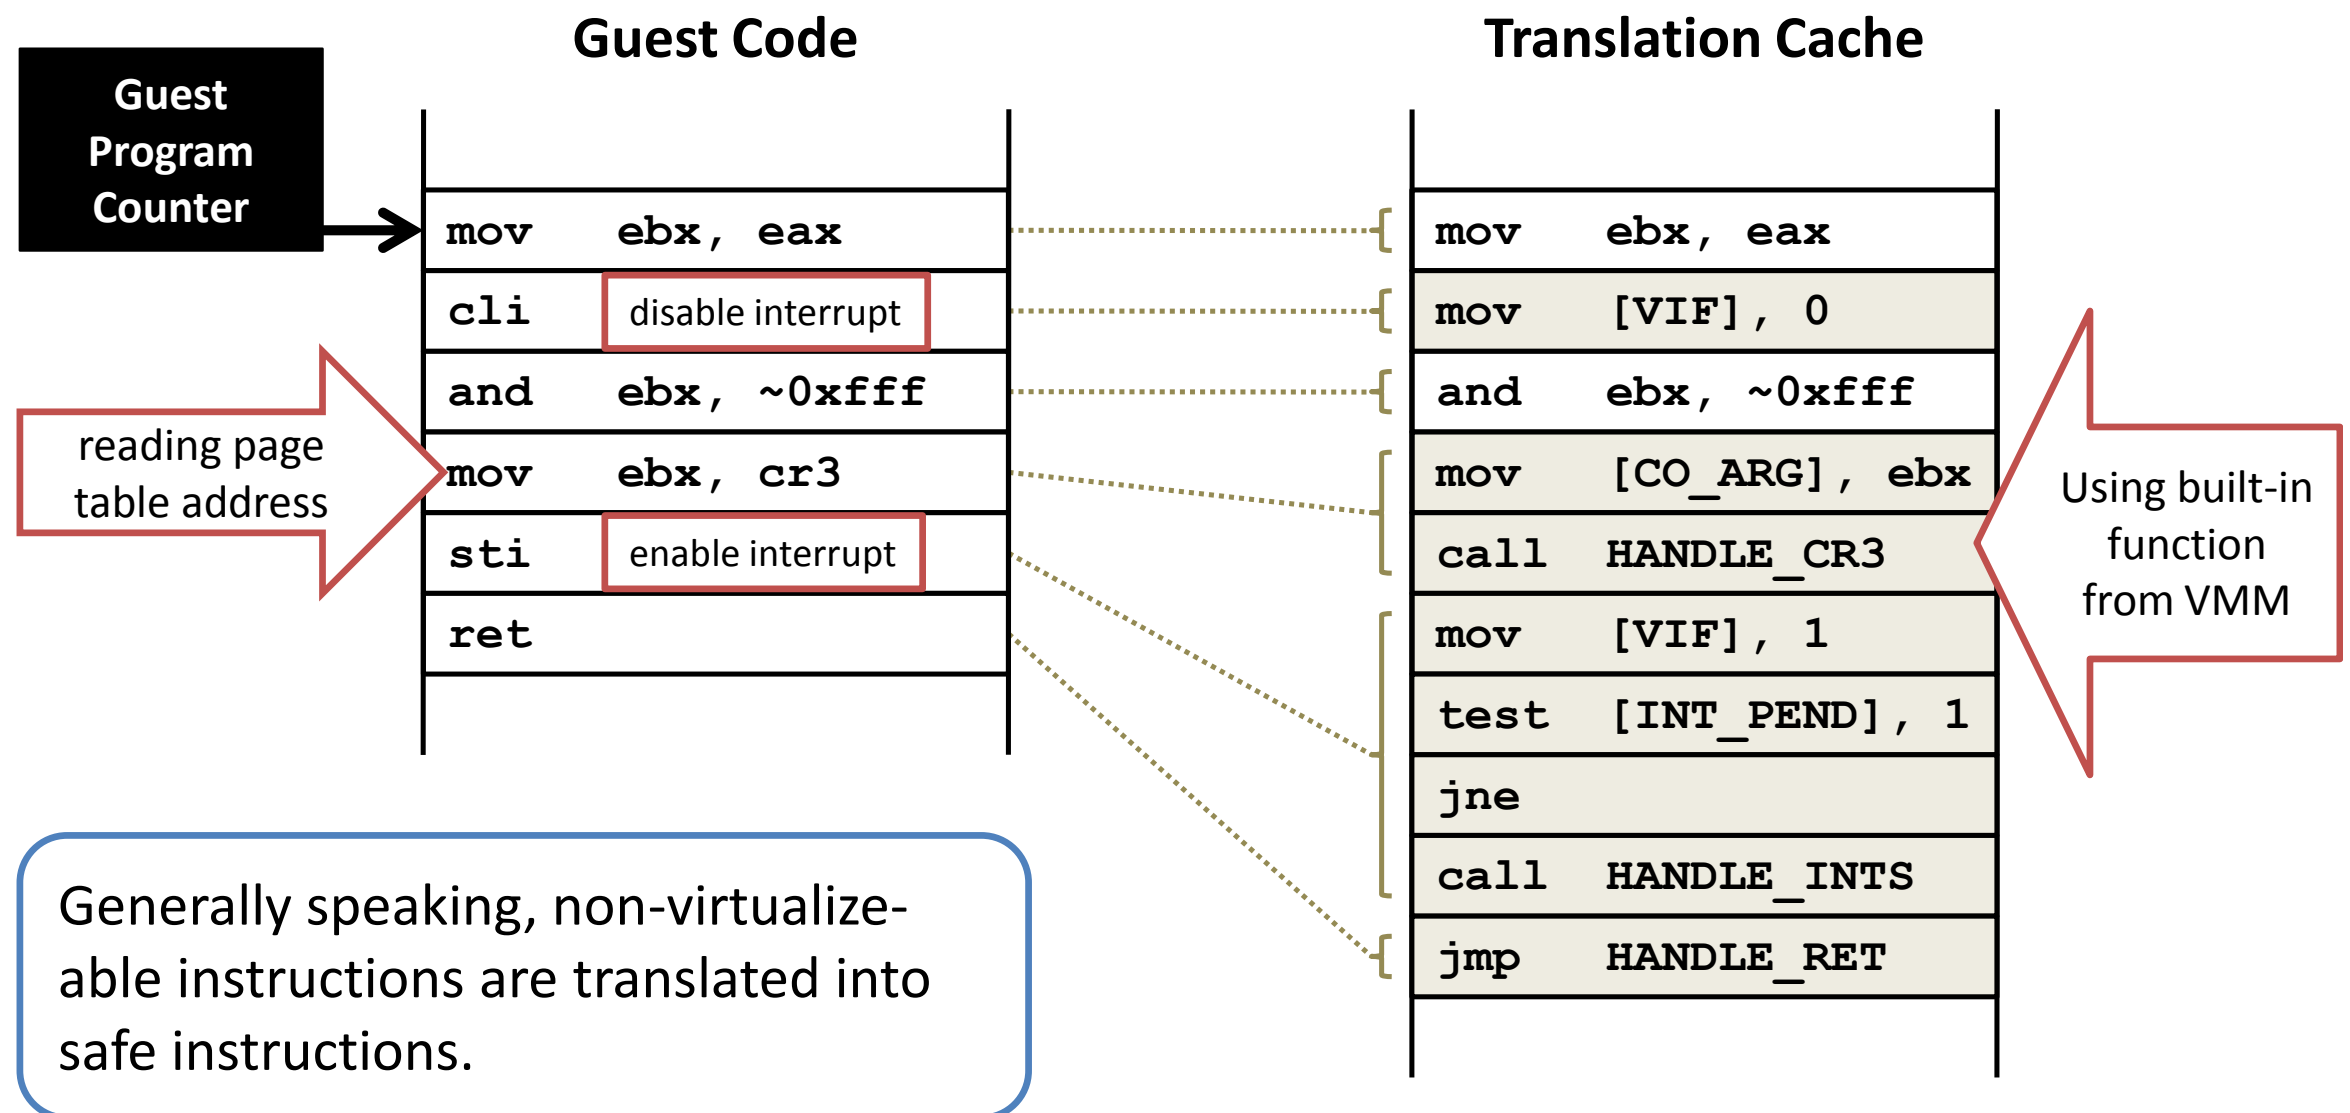
- Binary translation – step 1: trapping I/O calls

Ring 3	Guest applications
Ring 2	
Ring 1	Guest OS kernel
Ring 0	Hypervisor, Host OS
Host hardware	

whenever the guest OS asks for hardware, e.g. asking BIOS for a list of hardware, it's trapped by the hypervisor

Full virtualization

- Binary translation – step 2: emulate/translate



Full virtualization

- ▶ The guest OS is tricked to think that it's running privileged code in Ring 0, while it's actually running in Ring 1 of the host with the hypervisor emulating the hardware and trapping privileged code
- ▶ Unprivileged instructions are directly executed on CPU

Full virtualization

- ▶ Advantages:

- ▶ Keeps the guest OS unmodified
- ▶ Prevents an unstable VMs from impacting system performance; VM portability

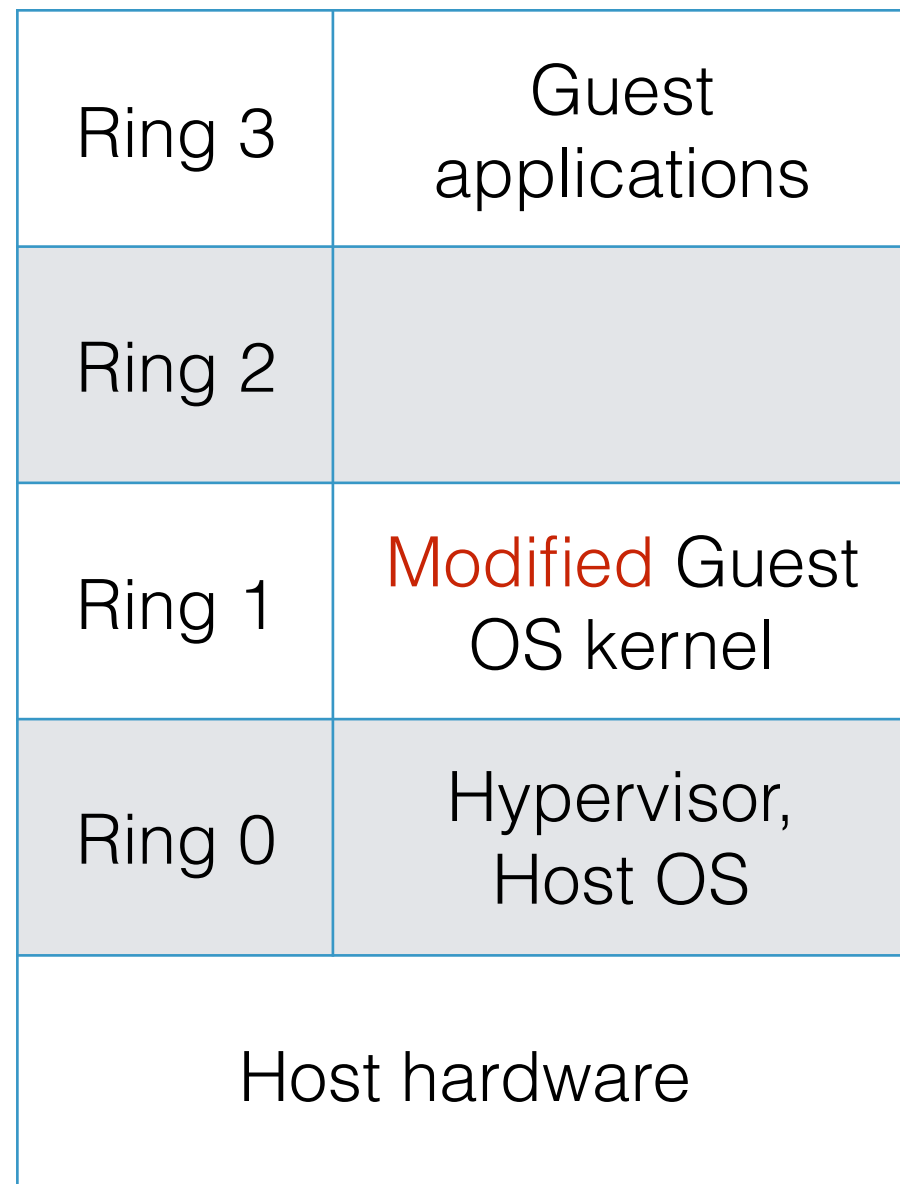
- ▶ Disadvantages:

- ▶ Performance is not good

Para-virtualization

- ▶ Developed to overcome the performance penalty of full virtualization with hardware emulation
- ▶ “Para” means “besides,” “with,” or “alongside.”

Para-virtualization



include virtualization
APIs and drivers

no binary translation

Para-virtualization

- ▶ Can be done in two ways:
 - ▶ A recompiled OS kernel. Easy for Linux, Windows doesn't support
 - ▶ Paravirtualization drivers for some hardware, e.g. GPU, NIC

Para-virtualization

- ▶ Guest OS is aware that it runs in a virtualized environment. It talks to the hypervisor through specialized APIs to run privileged instructions.
- ▶ These system calls, in the guest OS, are also called “hypercalls.”
- ▶ Performance is improved. The hypervisor can focus on isolating VMs and coordinating.

Hardware-assisted

Non-root mode	Ring 3	Guest applications
	Ring 2	
	Ring 1	
	Ring 0	Guest OS kernel
Root mode	Ring -1	Hypervisor
Host hardware		

likely to emerge as
the standard for
server virtualization
into the future

e.g., Intel® VT-x, AMD® V

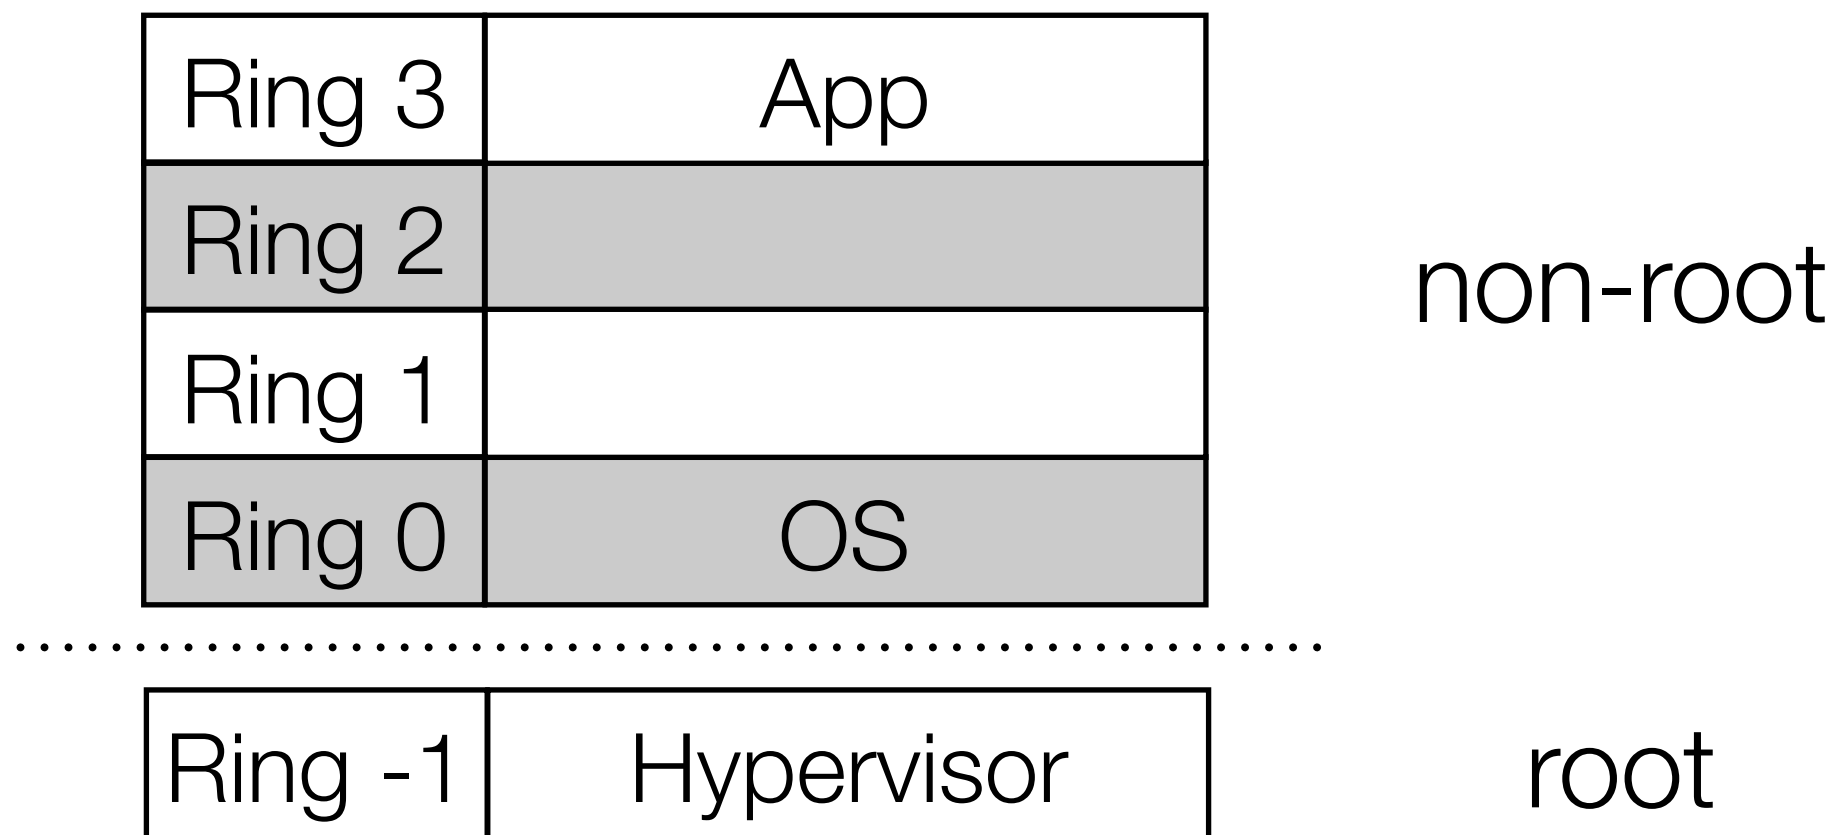
Hardware-assisted

Originally the machine is executing normally, without any guest OS.

Ring 3	App
Ring 2	
Ring 1	
Ring 0	OS

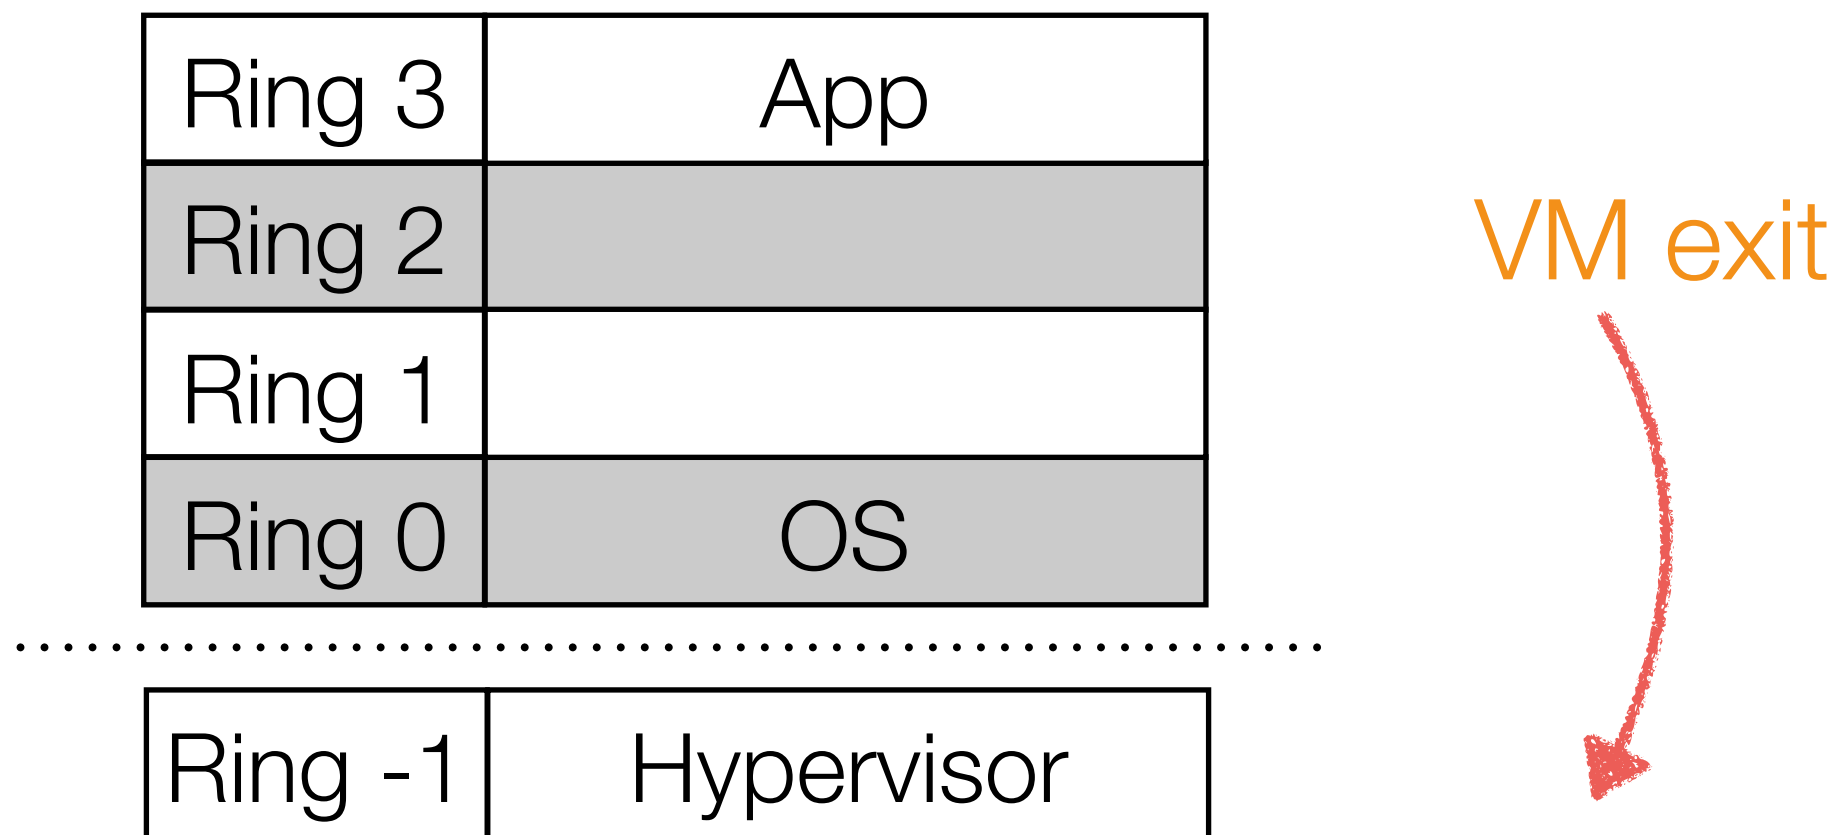
Hardware-assisted

When the hypervisor launches a VM,



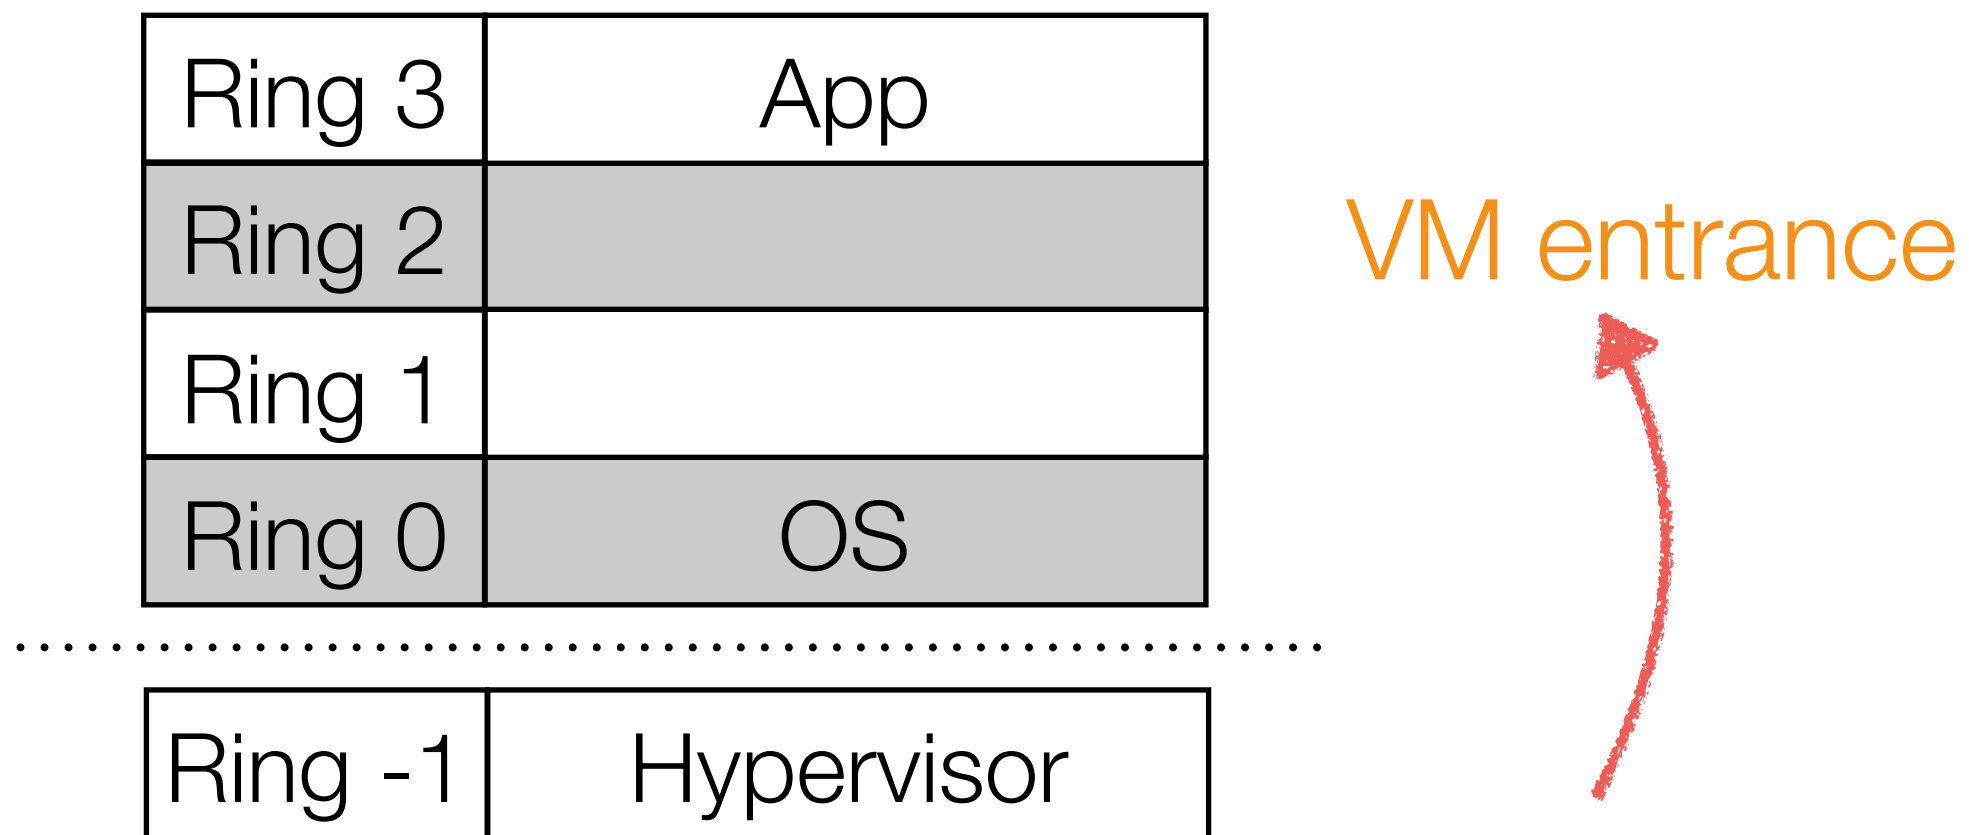
Hardware-assisted

When the guest OS meets certain triggers, which requires the hypervisor to exercise system control, a transition of control happens:




Hardware-assisted

When the hypervisor finishes, the control switches back to non-root mode, the VM continues



Hardware-assisted

The bigger picture here:



Each VM is allocated a specific hardware address space for performance isolation.

A Summary

	Full	Para-	Hardware-assisted
Handling privileged instructions	binary translation	hypercalls	non-root/ root mode
Guest OS modifications	no	yes	no
Performance	good	best	good
Examples	VMware, VirtualBox	Xen	Xen, VMware, VirtualBox, KVM

Credit

- ▶ Dr. WONG Tsz Yeung, CSCI 4180, CUHK
- ▶ Virtualization technology introduction, Intel Corporation.
28 July 2008.
- ▶ Wely Lau, wely@ncs.com.sg